

v27_12_v2

December 27, 2025

```
[61]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import roc_auc_score, f1_score, accuracy_score,
    average_precision_score, confusion_matrix, classification_report
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

# =====
# Project: Next-year financial distress prediction (firm-year panel)
# Data Science Lifecycle focus in this cell:
#   - Setup & configuration (temporal split conventions, leakage control)
#   - parameters)
#   - Panel integrity (types, deduplication, stable firm identifier)
#   - Define label_year used consistently for splitting and target alignment
# =====

# -----
# Configuration (split & preprocessing parameters)
# -----
FILE_NAME = "data.csv"

TRAIN_CUTOFF_LABEL_YEAR = 2022      # label_year <= cutoff → train/val pool; after
    ↪cutoff → test
VAL_YEARS = 1                         # last N years within the pool are validation
N_SPLITS_TIME_CV = 4                  # rolling time-based folds for sanity checks

WINSOR_LOWER_Q = 0.01                 # winsorization lower quantile (train-only)
WINSOR_UPPER_Q = 0.99                 # winsorization upper quantile (train-only)

REQUIRED_KEYS = ["gvkey", "fyear"]

# -----
# Utilities (robust numeric ops for ratios)
# -----
```

```

def to_float_numpy(x) -> np.ndarray:
    """Convert series/array-like to float numpy array, coercing non-numeric to
    ↪NaN."""
    s = pd.to_numeric(x, errors="coerce")
    return s.to_numpy(dtype=float) if hasattr(s, "to_numpy") else np.asarray(s, dtype=float)
def safe_divide(a, b) -> np.ndarray:
    """Elementwise divide a/b with NaN when division is invalid (0 or
    ↪non-finite)."""
    a = to_float_numpy(a)
    b = to_float_numpy(b)
    out = np.full_like(a, np.nan, dtype=float)
    np.divide(a, b, out=out, where=(b != 0) & np.isfinite(a) & np.isfinite(b))
    return out

def rolling_year_folds(
    df_in: pd.DataFrame, year_col: str = "label_year", n_splits: int = 4,
    ↪min_train_years: int = 3
) -> list[tuple[np.ndarray, np.ndarray, np.ndarray, int]]:
    """
    Create expanding-window time folds:
        train years: first (min_train_years + k) years
        val year:    next year
    Returns: list of (train_idx, val_idx, train_years, val_year)
    """
    years_sorted = np.sort(df_in[year_col].dropna().unique())
    if len(years_sorted) <= min_train_years:
        return []
    n_splits = min(n_splits, len(years_sorted) - min_train_years)

    folds_out = []
    for k in range(n_splits):
        train_years = years_sorted[: min_train_years + k]
        val_year = int(years_sorted[min_train_years + k])

        train_idx = df_in.index[df_in[year_col].isin(train_years)].to_numpy()
        val_idx = df_in.index[df_in[year_col] == val_year].to_numpy()
        folds_out.append((train_idx, val_idx, train_years, val_year))

    return folds_out

# =====
# 1) Data acquisition & panel hygiene
# =====
df = pd.read_csv(FILE_NAME, low_memory=False)

```

```

# Convert datadate if present
if "datadate" in df.columns:
    df["datadate"] = pd.to_datetime(df["datadate"], errors="coerce")

# Create stable firm id + de-duplicate firm-year (keep last record)
df["firm_id"] = df["gvkey"]
df = (
    df.sort_values(["firm_id", "fyear"])
    .drop_duplicates(subset=["firm_id", "fyear"], keep="last")
    .reset_index(drop=True)
)

# Label year: predict distress in the next fiscal year
df["label_year"] = df["fyear"] + 1

# =====
# 2) Scaffolding (define train/val pool years via label_year)
# =====
pool_mask = df["label_year"] <= TRAIN_CUTOFF_LABEL_YEAR
pool_years = np.sort(df.loc[pool_mask, "label_year"].dropna().unique())
val_years = pool_years[-VAL_YEARS:] if len(pool_years) else np.array([], dtype=int)

# This mask is ONLY used for imputations (train-only information)
train_mask_for_imputation = pool_mask & (~df["label_year"].isin(val_years))

```

[62]:

```

# =====
# Data Cleaning & Missing-Data Handling (leakage-aware)
# Purpose:
#   - Quantify missingness and distribution properties before intervention
#   - Preserve informative missingness via miss_* indicators
#   - Impute financial statement inputs using TRAIN-only information:
#     (1) within-firm past values (ffill) where economically sensible
#     (2) peer medians by year on size-scaled ratios
#   - Re-run EDA after imputation to audit how strongly imputations alter the data
# =====

RAW_INPUTS_FOR_FE = [
    "aco", "act", "ao", "aoloch", "ap", "apalch", "aqc", "at", "caps", "capx", "ceq", "che", "check", "csho",
    "datadate", "dlc", "dlcch", "dltis", "dltr", "dltt", "do", "dp", "dpc", "dv", "dvc", "dvp", "dvt", "esub",
    "fiao", "fincf", "fopo", "fyear", "gvkey", "ib", "ibadj", "ibc", "intan", "invch", "invt", "ismod", "iv",
    "ivao", "ivch", "ivncf", "ivstch", "lco", "lct", "lt", "mibt", "mkvalt", "niadj", "nopi", "oancf", "oib"
]

```

```

    ↵"prcc_c", "prcc_f", "prstkc", "pstk", "pstkn", "pstkr", "re", "recch", "rect", "seq", "siv", "spi", "sp
    ↵"sstk", "tstk", "txach", "txbcf", "txdc", "txditc", "txp", "txt", "xi", "xido", "xidoc", "xint",
      # optional identifiers present in many extracts:
      "conm", "consol", "datafmt", "indfmt",
]
raw = [c for c in RAW_INPUTS_FOR_FE if c in df.columns]

# -----
# 3.0 Ensure keys exist + types
# -----
# firm_id used for panel operations
if "firm_id" not in df.columns:
    if "gvkey" in df.columns:
        df["firm_id"] = df["gvkey"]
    else:
        raise ValueError("Need either firm_id or gvkey in df to run panel_
↪imputations.")

# fyear numeric
if "fyear" in df.columns:
    df["fyear"] = pd.to_numeric(df["fyear"], errors="coerce")

# datadate datetime (if present)
if "datadate" in df.columns:
    df["datadate"] = pd.to_datetime(df["datadate"], errors="coerce")

# -----
# 3.1 Drop rows with missing critical identifiers (do not impute these)
# -----
NON_IMPUTE_DROP = [c for c in ["gvkey", "datadate", "fyear", "conm", "datafmt", "indfmt", "consol"] if c in df.columns]
if NON_IMPUTE_DROP:
    before_n = df.shape[0]
    df = df.dropna(subset=NON_IMPUTE_DROP).copy()
    after_n = df.shape[0]
    if after_n < before_n:
        print(f"[INFO] Dropped {before_n - after_n:,} rows due to missing_
↪non-imputable ID/meta fields: {NON_IMPUTE_DROP}")

# Rebuild raw after potential drop
raw = [c for c in RAW_INPUTS_FOR_FE if c in df.columns]

# -----
# 3.2 EDA BEFORE imputation (missingness + distribution snapshot)
# -----

```

```

df_raw_pre = df[raw].copy(deep=True)

pre_miss = pd.DataFrame(
    {
        "col": raw,
        "n": [int(df_raw_pre[c].shape[0]) for c in raw],
        "n_na_pre": [int(df_raw_pre[c].isna().sum()) for c in raw],
        "pct_na_pre": [float(df_raw_pre[c].isna().mean() * 100.0) for c in raw],
        "train_n": [int(train_mask_for_imputation.sum()) for _ in raw],
        "train_pct_na_pre": [
            float(df_raw_pre.loc[train_mask_for_imputation, c].isna().mean() * 100.0) for c in raw
        ],
    },
)
.sort_values("pct_na_pre", ascending=False)

print("\n==== EDA (BEFORE imputation): Missingness on raw inputs ===")
print(pre_miss.round(4).head(50))

# Numeric distribution summary (exclude obvious non-numeric)
if raw:
    x_pre = df_raw_pre[raw].apply(pd.to_numeric, errors="coerce").replace([np.inf, -np.inf], np.nan)
    q_pre = x_pre.quantile([0.01, 0.05, 0.25, 0.50, 0.75, 0.95, 0.99]).T
    pre_dist = pd.DataFrame(
        {
            "n_nonmiss_pre": x_pre.notna().sum(),
            "mean_pre": x_pre.mean(),
            "std_pre": x_pre.std(ddof=0),
            "min_pre": x_pre.min(),
            "p01_pre": q_pre[0.01],
            "p05_pre": q_pre[0.05],
            "p25_pre": q_pre[0.25],
            "p50_pre": q_pre[0.50],
            "p75_pre": q_pre[0.75],
            "p95_pre": q_pre[0.95],
            "p99_pre": q_pre[0.99],
            "max_pre": x_pre.max(),
        }
    )
    print("\n==== EDA (BEFORE imputation): Distribution summary (raw inputs)====")
    print(pre_dist.round(4).sort_values("n_nonmiss_pre", ascending=True).head(50))

#
# -----
# 3.3 Missingness flags (ALWAYS create before imputations)

```

```

# -----
for c in raw:
    df[f"miss_{c}"] = df[c].isna().astype("int8")

# -----
# 3.4 Grouping key for peer-based imputations
# -----
group_cols = ["fyear"]

# 3.5 Step 1: Construct / reconcile FIRST (no leakage: uses contemporaneous or
#   ↪ lag only)
# -----
df = df.sort_values(["firm_id", "fyear"]).copy()

# 3.5.1 mkvalt construction from prcc_f * csho (if mkvalt missing)
if all(c in df.columns for c in ["mkvalt", "prcc_f", "csho"]):
    mkvalt_miss = df["mkvalt"].isna()
    mkvalt_calc = pd.to_numeric(df["prcc_f"], errors="coerce") * pd.
    ↪to_numeric(df["csho"], errors="coerce")
    df.loc[mkvalt_miss & mkvalt_calc.notna(), "mkvalt"] = mkvalt_calc.
    ↪loc[mkvalt_miss & mkvalt_calc.notna()]

# 3.5.2 Reconstruct change variables from level differences (fill only if
#   ↪ change var is missing)
def _fill_change_from_levels(change_col, level_col):
    if change_col in df.columns and level_col in df.columns:
        miss = df[change_col].isna()
        lvl = pd.to_numeric(df[level_col], errors="coerce")
        lag_lvl = df.groupby("firm_id")[level_col].shift(1)
        lag_lvl = pd.to_numeric(lag_lvl, errors="coerce")
        recon = lvl - lag_lvl
        df.loc[miss & recon.notna(), change_col] = recon.loc[miss & recon.
        ↪notna()]

_fill_change_from_levels("dlcch", "dlc")
_fill_change_from_levels("recch", "rect")
_fill_change_from_levels("invch", "invt")
_fill_change_from_levels("chech", "che")

# 3.5.3 apalch proxy with Δap if apalch missing
if "apalch" in df.columns and "ap" in df.columns:
    miss = df["apalch"].isna()
    ap = pd.to_numeric(df["ap"], errors="coerce")
    lag_ap = pd.to_numeric(df.groupby("firm_id")["ap"].shift(1), ↪
    ↪errors="coerce")
    recon = ap - lag_ap

```

```

df.loc[miss & recon.notna(), "apalch"] = recon.loc[miss & recon.notna()]

# -----
# 3.6 Step 2: ismod mode imputation (binary flag)
# -----
if "ismod" in df.columns:
    tr_obs = df.loc[train_mask_for_imputation, "ismod"]
    tr_obs_num = pd.to_numeric(tr_obs, errors="coerce")
    # global mode on training
    if tr_obs_num.notna().any():
        mode_val = float(tr_obs_num.mode(dropna=True).iloc[0])
    else:
        mode_val = 0.0
    df.loc[df["ismod"].isna(), "ismod"] = mode_val

# -----
# 3.7 Step 3: Stocks - firm-history (ffill) -> group median of ratio (x/at)
# -----
STOCKS = [c for c in [
    ↳ "aco", "act", "ao", "ap", "at", "caps", "ceq", "che", "csho", "cstk", "dlc", "dltt", "intan", "inv", "lco",
    ↳ "mibt", "ppent", "pstk", "pstkn", "pstkr", "re", "rect", "seq", "tstk", "ivaeq", "mkvalt"
] if c in df.columns]

# firm-history forward fill within each firm (uses only past observations)
if STOCKS:
    df[STOCKS] = df.groupby("firm_id")[STOCKS].ffill()

# ratio median impute remaining missing stocks (fit on training only)
NONNEG_STOCKS = set([c for c in STOCKS if c in [
    ↳ {"aco", "act", "ao", "ap", "at", "caps", "ceq", "che", "csho", "cstk", "dlc", "dltt",
    ↳ "intan", "inv", "lco", "lct", "lt", "mibt", "mkvalt", "ppent", "pstk", "pstkn", "pstkr",
    "rect", "seq", "tstk", "ivaeq"}])
])

def _fit_ratio_medians(train_df, col, size_col="at"):
    # returns (overall_ratio_median, series indexed by group tuple)
    s = pd.to_numeric(train_df[col], errors="coerce")
    size = pd.to_numeric(train_df[size_col], errors="coerce") if size_col in train_df.columns else None
    if size is None:
        # fallback: level median by groups
        grp_med = train_df.groupby(group_cols)[col].median()
        overall = float(s.median()) if s.notna().any() else 0.0
        return ("level", overall, grp_med)

```

```

valid = s.notna() & size.notna() & (size > 0)
if valid.sum() == 0:
    grp_med = train_df.groupby(group_cols)[col].median()
    overall = float(s.median()) if s.notna().any() else 0.0
    return ("level", overall, grp_med)

ratio = (s[valid] / size[valid]).replace([np.inf, -np.inf], np.nan).dropna()
overall = float(ratio.median()) if ratio.notna().any() else 0.0
tmp = train_df.loc[valid, group_cols].copy()
tmp["_ratio_"] = ratio.values
grp_med = tmp.groupby(group_cols)["_ratio_"].median()
return ("ratio", overall, grp_med)

def _apply_ratio_medians(df_all, col, fit_obj, size_col="at", nonneg=False):
    kind, overall, grp_med = fit_obj
    miss = df_all[col].isna()
    if not miss.any():
        return

    if kind == "ratio" and size_col in df_all.columns:
        size = pd.to_numeric(df_all.loc[miss, size_col], errors="coerce")
        # map group median ratio
        g = df_all.loc[miss, group_cols]
        mapped = pd.Series([np.nan] * miss.sum(), index=df_all.index[miss], ↴
                           dtype="float64")
        # build tuple key if multi-col grouping
        if len(group_cols) == 1:
            mapped = g[group_cols[0]].map(grp_med)
        else:
            keys = list(map(tuple, g[group_cols].to_numpy()))
            mapped = pd.Series(keys, index=g.index).map(grp_med)

        r = pd.to_numeric(mapped, errors="coerce").fillna(overall)
        fill = r * size
        # if size missing/≤0, fallback to level median within groups
        fill = fill.where(size.notna() & (size > 0), np.nan)
        df_all.loc[miss & fill.notna(), col] = fill.loc[miss & fill.notna()].to_numpy()

    # fallback: level median by group (still fit on training)
    # (use training group median of levels if ratio path didn't fill everything)
    miss2 = df_all[col].isna()
    if miss2.any():
        tr = df_all.loc[train_mask_for_imputation, [*group_cols, col]].copy()
        tr[col] = pd.to_numeric(tr[col], errors="coerce")
        lvl_overall = float(tr[col].median()) if tr[col].notna().any() else 0.0
        lvl_grp = tr.groupby(group_cols)[col].median()

```

```

g2 = df_all.loc[miss2, group_cols]
if len(group_cols) == 1:
    mapped2 = g2[group_cols[0]].map(lvl_grp)
else:
    keys2 = list(map(tuple, g2[group_cols].to_numpy()))
    mapped2 = pd.Series(keys2, index=g2.index).map(lvl_grp)

fill2 = pd.to_numeric(mapped2, errors="coerce").fillna(lvl_overall)
if nonneg:
    fill2 = fill2.clip(lower=0.0)
df_all.loc[miss2, col] = fill2.to_numpy()

# Fit on training
tr_all = df.loc[train_mask_for_imputation].copy()
ratio_fits = {}
for c in STOCKS:
    ratio_fits[c] = _fit_ratio_medians(tr_all, c, size_col="at")

# Apply to full df
for c in STOCKS:
    _apply_ratio_medians(df, c, ratio_fits[c], size_col="at", nonneg=(c in
NONNEG_STOCKS))

# -----
# 3.8 Step 4: Flows / income variables - ratio-median imputation (leakage-aware)
# -----
# Practical adjustments for this project:
# (i) Interest expense (xint) scales more naturally with debt than with
# assets.
#      => prefer imputing xint via (xint / total_debt) when debt is available.
# (ii) Tax components used in the FFO proxy (txt, txdc, txach) should not
#      block feature construction.
#      => set missing components to 0.0 (retain miss_* flags as indicators).
# (iii) Ratio-size imputations can explode for very large firms.
#      => clip only-imputed values to a robust training-observed quantile
# band.

FLOWS = [c for c in [
    "ib", "ibadj", "ibc", "niadj", "nopi", "oibdp", "dp", "txt", "oancf", "fincf", "ivncf", "xint", "esubc"
] if c in df.columns]

# --- Debt-aware base for xint (after stock imputations have populated dlc/
# dltt) ---
if "xint" in df.columns and all(c in df.columns for c in ["dlc", "dltt"]):

```

```

df["_td_for_xint"] = (
    pd.to_numeric(df["dlc"], errors="coerce").fillna(0.0)
    + pd.to_numeric(df["dltt"], errors="coerce").fillna(0.0)
)
else:
    df["_td_for_xint"] = np.nan

# Debt-aware rule: if total debt == 0, interest expense is set to 0 (keep missu
# →flag already created)
if "xint" in df.columns and "_td_for_xint" in df.columns:
    xint_miss = df["xint"].isna()
    df.loc[xint_miss & (pd.to_numeric(df["_td_for_xint"], errors="coerce").
    fillna(0.0) <= 0), "xint"] = 0.0

# Fit ratio medians on training and apply (xint uses debt base when possible;u
# →otherwise assets base)
flow.fits = {}
for c in FLOWS:
    size_base = "_td_for_xint" if (c == "xint" and "_td_for_xint" in df.
    columns) else "at"
    flow.fits[c] = _fit_ratio_medians(tr_all, c, size_col=size_base)

for c in FLOWS:
    size_base = "_td_for_xint" if (c == "xint" and "_td_for_xint" in df.
    columns) else "at"
    # flows can be negative; do not clip to non-negative
    _apply_ratio_medians(df, c, flow.fits[c], size_col=size_base, nonneg=False)

# --- Practical fill: tax components used in FFO proxy ---
for c in ["txt", "txdc", "txach"]:
    if c in df.columns:
        df.loc[df[c].isna(), c] = 0.0

# --- Practical guardrail: cap only-imputed values to training-observedu
# →quantiles ---
def _cap_imputed_to_train_quantiles(
    df_all: pd.DataFrame,
    col: str,
    lower_q: float = 0.01,
    upper_q: float = 0.99,
    nonneg: bool = False,
) -> None:
    flag = f"miss_{col}"
    if col not in df_all.columns or flag not in df_all.columns:
        return

```

```

miss_mask = df_all[flag].astype(bool)
if not miss_mask.any():
    return

# Prefer bounds from observed (not originally missing) values in the
# training split
obs = pd.to_numeric(df_all.loc[train_mask_for_imputation & (~miss_mask), col], errors="coerce")
if obs.notna().sum() < 200:
    obs = pd.to_numeric(df_all.loc[train_mask_for_imputation, col], errors="coerce")

if obs.notna().sum() == 0:
    return

lo = float(obs.quantile(lower_q))
hi = float(obs.quantile(upper_q))

if nonneg:
    lo = max(0.0, lo) if np.isfinite(lo) else 0.0

s = pd.to_numeric(df_all.loc[miss_mask, col], errors="coerce")
df_all.loc[miss_mask, col] = s.clip(lo, hi).to_numpy()

# Apply capping to the variables that are later used (directly or indirectly)
# in engineered ratios
CAP_COLS = sorted(set(STOCKS + FLOWS + ["dlcch", "apalch", "recch", "invch",
                                         "chech"]) & set(df.columns))
for c in CAP_COLS:
    _cap_imputed_to_train_quantiles(df, c, lower_q=0.01, upper_q=0.99,
                                     nonneg=(c in NONNEG_STOCKS))

# 3.9 EDA AFTER imputation (missingness reduction + distribution deltas)
# -----
df_raw_post = df[raw].copy(deep=True)

post_miss = pd.DataFrame(
{
    "col": raw,
    "n_na_post": [int(df_raw_post[c].isna().sum()) for c in raw],
    "pct_na_post": [float(df_raw_post[c].isna().mean() * 100.0) for c in raw],
    "train_pct_na_post": [
        float(df_raw_post.loc[train_mask_for_imputation, c].isna().mean() * 100.0) for c in raw
    ]
})

```

```

        ],
    }
)

changes = pre_miss.merge(post_miss, on="col", how="left")
changes["n_imputed"] = changes["n_na_pre"] - changes["n_na_post"]
changes["pct_points_na_reduction"] = changes["pct_na_pre"] - □
    ↵ changes["pct_na_post"]

x_post = df_raw_post[raw].apply(pd.to_numeric, errors="coerce").replace([np.
    ↵ inf, -np.inf], np.nan)
q_post = x_post.quantile([0.01, 0.05, 0.25, 0.50, 0.75, 0.95, 0.99]).T
post_dist = pd.DataFrame(
{
    "n_nonmiss_post": x_post.notna().sum(),
    "mean_post": x_post.mean(),
    "std_post": x_post.std(ddof=0),
    "min_post": x_post.min(),
    "p01_post": q_post[0.01],
    "p05_post": q_post[0.05],
    "p25_post": q_post[0.25],
    "p50_post": q_post[0.50],
    "p75_post": q_post[0.75],
    "p95_post": q_post[0.95],
    "p99_post": q_post[0.99],
    "max_post": x_post.max(),
}
)
)

pre_dist_key = □
    ↵ pre_dist[["n_nonmiss_pre", "mean_pre", "std_pre", "p01_pre", "p50_pre", "p99_pre"]].
    ↵ copy() if raw else pd.DataFrame()
post_dist_key = □
    ↵ post_dist[["n_nonmiss_post", "mean_post", "std_post", "p01_post", "p50_post", "p99_post"]].
    ↵ copy() if raw else pd.DataFrame()

dist_delta = pre_dist_key.join(post_dist_key, how="outer")
dist_delta["delta_mean"] = dist_delta["mean_post"] - dist_delta["mean_pre"]
dist_delta["delta_std"] = dist_delta["std_post"] - dist_delta["std_pre"]
dist_delta["delta_p50"] = dist_delta["p50_post"] - dist_delta["p50_pre"]

# Imputed-only diagnostics
rows = []
for c in raw:
    pre_na_mask = df_raw_pre[c].isna()
    n_imp = int(pre_na_mask.sum())
    if n_imp == 0:

```

```

        rows.append((c, 0, np.nan, np.nan, np.nan, np.nan, np.nan))
    continue

    imp_vals = pd.to_numeric(df_raw_post.loc[pre_na_mask, c], errors="coerce").
    ↪replace([np.inf, -np.inf], np.nan)
    obs_vals = pd.to_numeric(df_raw_pre.loc[~pre_na_mask, c], errors="coerce").
    ↪replace([np.inf, -np.inf], np.nan)

    rows.append(
        (
            c,
            n_imp,
            float(imp_vals.mean()) if imp_vals.notna().any() else np.nan,
            float(imp_vals.median()) if imp_vals.notna().any() else np.nan,
            float(imp_vals.std(ddof=0)) if imp_vals.notna().any() else np.nan,
            float(obs_vals.mean()) if obs_vals.notna().any() else np.nan,
            float(obs_vals.median()) if obs_vals.notna().any() else np.nan,
        )
    )

imputed_only = pd.DataFrame(
    rows,
    ↳
    ↪columns=["col", "n_imputed", "imputed_mean", "imputed_median", "imputed_std", "observed_mean_pre"]
).set_index("col")

print("\n==== EDA (AFTER imputation): Missingness on raw inputs + change ===")
cols_show = [
    "col", "n", "n_na_pre", "pct_na_pre", "n_na_post", "pct_na_post",
    "n_imputed", "pct_points_na_reduction", "train_pct_na_pre", ↳
    ↪"train_pct_na_post",
]
print(
    changes[cols_show]
    .sort_values(["n_imputed", "pct_points_na_reduction"], ascending=[False, ↳
    ↪False])
    .round(4)
    .head(50)
)

print("\n==== Change analysis: Distribution deltas (post - pre) on raw inputs ↳
    ↪====")
print(
    ↳
    ↪dist_delta[["n_nonmiss_pre", "n_nonmiss_post", "delta_mean", "delta_std", "delta_p50"]]
    .sort_values("delta_mean", key=lambda s: s.abs(), ascending=False)
)

```

```

    .round(6)
    .head(50)
)

print("\n==== Change analysis: Imputed-only vs observed (pre) summary ===")
print(
    imputed_only.assign(
        mean_gap_imputed_minus_observed=lambda d: d["imputed_mean"] - d["observed_mean_pre"],
        median_gap_imputed_minus_observed=lambda d: d["imputed_median"] - d["observed_median_pre"],
    )
    .sort_values("n_imputed", ascending=False)
    .round(6)
    .head(50)
)

```

==== EDA (BEFORE imputation): Missingness on raw inputs ===

	col	n	n_na_pre	pct_na_pre	train_n	train_pct_na_pre
18	dlcch	75005	33143	44.1877	48458	42.9630
5	apalch	75005	30371	40.4920	48458	39.0214
75	txach	75005	22791	30.3860	48458	29.2501
48	ivstch	75005	19194	25.5903	48458	23.0282
66	recch	75005	12589	16.7842	48458	16.5938
53	mkvalt	75005	12350	16.4656	48458	17.0189
71	sppe	75005	12239	16.3176	48458	16.4307
1	act	75005	10721	14.2937	48458	14.6581
50	lct	75005	10695	14.2590	48458	14.5982
84	xint	75005	10536	14.0471	48458	13.9234
78	txditc	75005	9069	12.0912	48458	12.2044
79	txp	75005	7942	10.5886	48458	10.6917
29	esubc	75005	6798	9.0634	48458	8.8964
49	lco	75005	6779	9.0381	48458	9.3277
0	aco	75005	6778	9.0367	48458	9.3277
60	prcc_f	75005	5956	7.9408	48458	8.8902
59	prcc_c	75005	5942	7.9221	48458	8.8819
40	invch	75005	5315	7.0862	48458	7.3445
72	sppiv	75005	3967	5.2890	48458	4.9631
44	ivaeq	75005	3548	4.7304	48458	4.8826
61	prstkc	75005	3410	4.5464	48458	4.9486
8	caps	75005	3074	4.0984	48458	4.2655
6	aqc	75005	2812	3.7491	48458	3.2647
23	dp	75005	2722	3.6291	48458	3.7909
65	re	75005	2596	3.4611	48458	3.6650
46	ivch	75005	2540	3.3864	48458	3.3844
45	ivao	75005	2529	3.3718	48458	3.1821

77	txdc	75005	2472	3.2958	48458	3.4050
69	siv	75005	2299	3.0651	48458	3.1285
57	oibdp	75005	2142	2.8558	48458	2.9366
14	cstk	75005	1925	2.5665	48458	3.0294
24	dpc	75005	1757	2.3425	48458	2.4268
19	dltis	75005	1756	2.3412	48458	2.4681
58	ppent	75005	1723	2.2972	48458	2.3814
73	sstk	75005	1436	1.9145	48458	1.8593
20	dltr	75005	1388	1.8505	48458	2.0616
25	dv	75005	1137	1.5159	48458	1.5911
70	spi	75005	1126	1.5012	48458	1.3785
67	rect	75005	880	1.1733	48458	1.1247
52	mibt	75005	842	1.1226	48458	1.3022
13	csho	75005	804	1.0719	48458	1.2939
28	dvt	75005	706	0.9413	48458	1.1412
26	dvc	75005	703	0.9373	48458	1.1350
39	intan	75005	654	0.8719	48458	0.8977
41	invt	75005	652	0.8693	48458	0.8255
76	txbcf	75005	619	0.8253	48458	1.0071
9	capx	75005	503	0.6706	48458	0.7429
83	xidoc	75005	477	0.6360	48458	0.7202
3	aoloch	75005	462	0.6160	48458	0.6893
38	ibc	75005	454	0.6053	48458	0.6810

==== EDA (BEFORE imputation): Distribution summary (raw inputs) ===

	n_nonmiss_pre	mean_pre	std_pre	min_pre	\
consol	0	NaN	NaN	NaN	
datafmt	0	NaN	NaN	NaN	
conm	0	NaN	NaN	NaN	
indfmt	0	NaN	NaN	NaN	
dlcch	41862	2.168121e+03	1.155250e+06	-8.555368e+07	
apalch	44634	8.385212e+03	3.253932e+05	-4.849894e+07	
txach	52214	3.214308e+02	3.140000e+04	-3.282293e+06	
ivstch	55811	6.871673e+04	3.980372e+06	-1.752327e+07	
recch	62416	-1.318528e+04	3.129434e+05	-2.297888e+07	
mkvalt	62655	9.288907e+05	2.201298e+07	0.000000e+00	
sppe	62766	6.939724e+03	1.245321e+05	-5.895100e+04	
act	64284	6.440956e+05	4.645197e+06	-2.498000e+03	
lct	64310	4.755088e+05	3.980895e+06	0.000000e+00	
xint	64469	2.759126e+04	1.196406e+05	-7.284730e+05	
txditz	65936	4.384545e+04	3.265589e+05	-1.616000e+03	
txp	67063	1.062620e+04	1.001099e+05	-2.217000e+03	
esubc	68207	-2.220041e+03	8.665624e+04	-8.775761e+06	
lco	68226	2.039805e+05	1.784063e+06	-8.900000e+01	
aco	68227	5.573244e+04	9.308496e+05	-2.500000e+01	
prcc_f	69029	6.838158e+02	1.018087e+04	1.000000e-04	
prcc_c	69039	7.067634e+02	1.003861e+04	1.000000e-04	
invch	69690	-3.860008e+04	1.008997e+06	-8.472324e+07	

sppiv	71038	-7.943608e+03	1.753101e+05	-1.598761e+07	
ivaeq	71457	6.873078e+04	9.162568e+05	-1.551500e+04	
prstkc	71595	2.902337e+04	2.408411e+05	-1.131100e+04	
caps	71931	5.224117e+05	2.924028e+06	-2.333792e+07	
aqc	72193	2.583727e+04	3.427732e+05	-2.002407e+07	
dp	72283	7.989028e+04	7.010127e+05	-4.121000e+03	
re	72409	4.063079e+05	4.773704e+06	-8.108442e+07	
ivch	72465	4.336588e+05	1.267386e+07	0.000000e+00	
ivao	72476	1.117387e+06	1.830844e+07	0.000000e+00	
txdc	72533	-2.552908e+03	8.582746e+04	-8.127405e+06	
siv	72706	3.204806e+05	9.796796e+06	-3.080000e+02	
oibdp	72863	2.472644e+05	1.847387e+06	-6.234977e+06	
cstk	73080	1.440901e+05	1.035999e+06	-1.466100e+04	
dpc	73248	8.009984e+04	6.495898e+05	-1.551370e+05	
dltis	73249	2.584506e+05	2.896443e+06	-4.257400e+04	
ppent	73282	5.870571e+05	5.242606e+06	0.000000e+00	
sstk	73569	2.949253e+04	2.324514e+05	-5.012830e+05	
dltr	73617	2.679871e+05	2.856270e+06	-2.844300e+04	
dv	73868	4.155940e+04	3.023818e+05	-1.631969e+06	
spi	73879	-1.290533e+04	1.992509e+05	-1.354613e+07	
rect	74125	2.121967e+06	3.060638e+07	-7.300000e+01	
mibt	74163	5.164076e+04	6.023643e+05	-1.237174e+06	
csho	74201	2.254364e+05	5.059555e+06	0.000000e+00	
dvt	74299	4.154428e+04	3.087856e+05	-1.329470e+05	
dvc	74302	4.081004e+04	3.071635e+05	-2.473000e+03	
intan	74351	3.830857e+05	2.603599e+06	0.000000e+00	
invt	74353	1.471455e+05	1.356403e+06	0.000000e+00	
txbcnf	74386	2.081901e+02	4.374733e+03	-9.510600e+04	

	p01_pre	p05_pre	p25_pre	p50_pre	p75_pre	\
consol	NaN	NaN	NaN	NaN	NaN	
datafmt	NaN	NaN	NaN	NaN	NaN	
conm	NaN	NaN	NaN	NaN	NaN	
indfmt	NaN	NaN	NaN	NaN	NaN	
dlcch	-2.265422e+05	-8701.6500	0.0000	0.0000	1.0000	
apalch	-1.053346e+05	-14648.8000	-119.0000	50.0000	1693.0000	
txach	-9.192830e+03	-193.0000	0.0000	0.0000	0.0000	
ivstch	-2.099582e+05	-10799.5000	0.0000	0.0000	0.0000	
recch	-2.890472e+05	-52494.2500	-2252.2500	-17.0000	54.0000	
mkvalt	5.154000e-01	3.3549	48.9974	467.1767	4760.4262	
sppe	0.000000e+00	0.0000	0.0000	0.0000	6.8300	
act	1.700000e-01	21.0000	1732.0000	23074.0000	214833.2500	
lct	6.600000e-01	26.5090	1277.0000	11405.0000	105630.0000	
xint	0.000000e+00	0.0000	15.0000	433.0000	8992.0000	
txditc	0.000000e+00	0.0000	0.0000	0.0000	1883.5000	
txp	0.000000e+00	0.0000	0.0000	0.0000	152.0000	
esubc	-3.866582e+04	-484.7000	0.0000	0.0000	0.0000	
lco	0.000000e+00	0.0000	301.0000	3986.5000	46136.5000	

aco	0.000000e+00	0.0000	58.0000	946.0000	10293.0000
prcc_f	1.000000e-02	0.1016	2.6000	13.9200	42.8400
prcc_c	1.000000e-02	0.1000	2.6200	14.0100	43.0250
invch	-3.362948e+05	-33105.1500	-172.0000	0.0000	0.0000
sppiv	-1.356993e+05	-11708.4500	-13.0000	0.0000	0.0000
ivaeq	0.000000e+00	0.0000	0.0000	0.0000	0.0000
prstkc	0.000000e+00	0.0000	0.0000	0.0000	846.5000
caps	0.000000e+00	0.0000	1645.0000	26381.0000	268051.0000
aqc	-1.760872e+04	0.0000	0.0000	0.0000	0.0000
dp	0.000000e+00	0.0000	68.0000	1079.4700	15628.5000
re	-2.484404e+06	-701966.6000	-80817.0000	-1251.0000	29686.0000
ivch	0.000000e+00	0.0000	0.0000	0.0000	275.0000
ivao	0.000000e+00	0.0000	0.0000	0.0000	3979.7500
txdc	-1.337068e+05	-18598.4000	-105.0000	0.0000	5.0000
siv	0.000000e+00	0.0000	0.0000	0.0000	61.0000
oibdp	-2.060061e+05	-47857.8000	-1873.5000	800.4000	46909.0000
cstk	0.000000e+00	0.0000	8.0000	197.0000	13957.2500
dpc	0.000000e+00	0.0000	100.3000	1542.0000	21382.2500
dltis	0.000000e+00	0.0000	0.0000	0.1500	1740.0000
ppent	0.000000e+00	0.0000	514.0000	8837.5000	93080.2500
sstk	0.000000e+00	0.0000	0.0000	58.0000	4216.0000
dltr	0.000000e+00	0.0000	0.0000	140.7000	8902.0000
dv	0.000000e+00	0.0000	0.0000	0.0000	884.0000
spi	-3.215187e+05	-47237.4000	-1071.0000	0.0000	0.0000
rect	0.000000e+00	0.0000	225.0000	5841.0000	97807.0000
mibt	-1.887380e+03	0.0000	0.0000	0.0000	0.0000
csho	3.150000e+00	28.0000	7219.0000	36145.0000	107343.0000
dvt	0.000000e+00	0.0000	0.0000	0.0000	1178.4450
dvc	0.000000e+00	0.0000	0.0000	0.0000	621.0000
intan	0.000000e+00	0.0000	0.0000	1365.0000	45975.0000
invt	0.000000e+00	0.0000	0.0000	275.3000	12041.0000
txbcf	0.000000e+00	0.0000	0.0000	0.0000	0.0000

	p95_pre	p99_pre	max_pre
consol	NaN	NaN	NaN
datafmt	NaN	NaN	NaN
conm	NaN	NaN	NaN
indfmt	NaN	NaN	NaN
dlcch	12288.900	239569.72	1.674825e+08
apalch	40626.050	280635.99	2.085900e+07
txach	544.000	16622.89	3.282293e+06
ivstch	14654.500	206862.80	3.964813e+08
recch	14636.250	102245.15	3.777953e+07
mkvalt	1137140.389	16710068.76	3.522211e+09
sppe	7315.750	96130.00	7.819213e+06
act	1887389.100	11141665.87	2.648894e+08
lct	1150175.850	8749902.35	2.275619e+08
xint	126670.400	475515.16	3.895731e+06

txditc	138738.750	824518.10	1.309682e+07
txp	20801.100	229018.76	5.843978e+06
esubc	503.700	18641.34	3.314734e+06
lco	528718.500	3510482.25	1.976466e+08
aco	124386.000	926127.42	1.973027e+08
prcc_f	271.900	17260.04	1.369125e+06
prcc_c	275.000	17524.72	1.369125e+06
invch	9558.700	99906.23	1.919409e+07
sppiv	779.000	12781.32	7.378669e+06
ivaeq	69436.200	1094948.44	4.184002e+07
prstkc	87371.300	608747.74	1.551860e+07
caps	2034165.500	7334828.10	1.724641e+08
aqc	71387.800	578083.72	5.188218e+07
dp	220915.500	1361138.02	3.714574e+07
re	1402781.800	11353646.12	2.219454e+08
ivch	441795.800	3486163.72	1.169666e+09
ivao	1193829.500	10658327.50	1.099816e+09
txdc	11204.800	85023.32	4.507005e+06
siv	237122.750	2260226.70	9.068858e+08
oibdp	727543.600	4683708.38	7.350011e+07
cstk	420270.900	2924964.10	3.904769e+07
dpc	252347.000	1229121.52	3.648387e+07
dltis	692136.800	4321096.08	1.818726e+08
ppent	1817488.900	10065753.57	2.926841e+08
sstk	111376.400	525775.68	2.036114e+07
dltr	761854.400	3911400.36	1.796217e+08
dv	122377.700	899427.56	1.296791e+07
spi	2226.100	54763.82	1.148318e+07
rect	1917288.600	17121476.16	1.218880e+09
mibt	66241.300	1067080.28	3.075527e+07
csho	533438.000	2530459.00	6.064077e+08
dvt	124549.500	879155.68	1.485872e+07
dvc	120572.700	869661.95	1.437446e+07
intan	1345701.500	7366060.00	9.849622e+07
invt	385661.800	2910600.48	1.043358e+08
txbcnf	0.000	3680.50	6.321680e+05

==== EDA (AFTER imputation): Missingness on raw inputs + change ===

	col	n	n_na_pre	pct_na_pre	n_na_post	pct_na_post	n_imputed	\
0	dlcch	75005	33143	44.1877	4502	6.0023	28641	
1	apalch	75005	30371	40.4920	4391	5.8543	25980	
2	txach	75005	22791	30.3860	0	0.0000	22791	
5	mkvalt	75005	12350	16.4656	0	0.0000	12350	
7	act	75005	10721	14.2937	0	0.0000	10721	
8	lct	75005	10695	14.2590	0	0.0000	10695	
9	xint	75005	10536	14.0471	0	0.0000	10536	
4	recch	75005	12589	16.7842	2473	3.2971	10116	
12	esubc	75005	6798	9.0634	0	0.0000	6798	

13	lco	75005	6779	9.0381	0	0.0000	6779
14	aco	75005	6778	9.0367	0	0.0000	6778
17	invch	75005	5315	7.0862	1241	1.6546	4074
19	ivaeq	75005	3548	4.7304	0	0.0000	3548
21	caps	75005	3074	4.0984	0	0.0000	3074
23	dp	75005	2722	3.6291	0	0.0000	2722
24	re	75005	2596	3.4611	0	0.0000	2596
27	txdc	75005	2472	3.2958	0	0.0000	2472
29	oibdp	75005	2142	2.8558	0	0.0000	2142
30	cstk	75005	1925	2.5665	0	0.0000	1925
33	ppent	75005	1723	2.2972	0	0.0000	1723
38	rect	75005	880	1.1733	0	0.0000	880
39	mibt	75005	842	1.1226	0	0.0000	842
40	csho	75005	804	1.0719	0	0.0000	804
43	intan	75005	654	0.8719	0	0.0000	654
44	invt	75005	652	0.8693	0	0.0000	652
49	ibc	75005	454	0.6053	0	0.0000	454
54	ivncf	75005	316	0.4213	0	0.0000	316
55	fincf	75005	315	0.4200	0	0.0000	315
57	oancf	75005	303	0.4040	0	0.0000	303
56	chech	75005	304	0.4053	57	0.0760	247
58	pstk	75005	219	0.2920	0	0.0000	219
59	pstkn	75005	211	0.2813	0	0.0000	211
60	dltt	75005	188	0.2506	0	0.0000	188
61	tstk	75005	180	0.2400	0	0.0000	180
62	ceq	75005	149	0.1987	0	0.0000	149
63	lt	75005	101	0.1347	0	0.0000	101
64	ap	75005	88	0.1173	0	0.0000	88
65	pstkr	75005	75	0.1000	0	0.0000	75
67	dlc	75005	31	0.0413	0	0.0000	31
68	txt	75005	15	0.0200	0	0.0000	15
69	nopi	75005	12	0.0160	0	0.0000	12
70	che	75005	6	0.0080	0	0.0000	6
71	ao	75005	5	0.0067	0	0.0000	5
74	seq	75005	3	0.0040	0	0.0000	3
3	ivstch	75005	19194	25.5903	19194	25.5903	0
6	sppe	75005	12239	16.3176	12239	16.3176	0
10	txditc	75005	9069	12.0912	9069	12.0912	0
11	txp	75005	7942	10.5886	7942	10.5886	0
15	prcc_f	75005	5956	7.9408	5956	7.9408	0
16	prcc_c	75005	5942	7.9221	5942	7.9221	0

	pct_points_na_reduction	train_pct_na_pre	train_pct_na_post
0	38.1855	42.9630	8.1844
1	34.6377	39.0214	7.8645
2	30.3860	29.2501	0.0000
5	16.4656	17.0189	0.0000
7	14.2937	14.6581	0.0000

8	14.2590	14.5982	0.0000
9	14.0471	13.9234	0.0000
4	13.4871	16.5938	4.0303
12	9.0634	8.8964	0.0000
13	9.0381	9.3277	0.0000
14	9.0367	9.3277	0.0000
17	5.4316	7.3445	2.0265
19	4.7304	4.8826	0.0000
21	4.0984	4.2655	0.0000
23	3.6291	3.7909	0.0000
24	3.4611	3.6650	0.0000
27	3.2958	3.4050	0.0000
29	2.8558	2.9366	0.0000
30	2.5665	3.0294	0.0000
33	2.2972	2.3814	0.0000
38	1.1733	1.1247	0.0000
39	1.1226	1.3022	0.0000
40	1.0719	1.2939	0.0000
43	0.8719	0.8977	0.0000
44	0.8693	0.8255	0.0000
49	0.6053	0.6810	0.0000
54	0.4213	0.4788	0.0000
55	0.4200	0.4767	0.0000
57	0.4040	0.4602	0.0000
56	0.3293	0.4623	0.1156
58	0.2920	0.2951	0.0000
59	0.2813	0.2868	0.0000
60	0.2506	0.2889	0.0000
61	0.2400	0.2414	0.0000
62	0.1987	0.2043	0.0000
63	0.1347	0.1259	0.0000
64	0.1173	0.0991	0.0000
65	0.1000	0.1032	0.0000
67	0.0413	0.0495	0.0000
68	0.0200	0.0144	0.0000
69	0.0160	0.0227	0.0000
70	0.0080	0.0124	0.0000
71	0.0067	0.0103	0.0000
74	0.0040	0.0062	0.0000
3	0.0000	23.0282	23.0282
6	0.0000	16.4307	16.4307
10	0.0000	12.2044	12.2044
11	0.0000	10.6917	10.6917
15	0.0000	8.8902	8.8902
16	0.0000	8.8819	8.8819

==== Change analysis: Distribution deltas (post - pre) on raw inputs ====
n_nonmiss_pre n_nonmiss_post delta_mean delta_std \

mkvalt	62655	75005	342407.151082	-1.766379e+06
act	64284	75005	185811.385530	-1.788945e+05
lct	64310	75005	90560.417655	-2.078001e+05
lco	68226	75005	25133.438729	-6.005804e+04
rect	74125	75005	-23898.875943	-1.792580e+05
re	72409	75005	-16713.635107	-8.185532e+04
lt	74904	75005	-6139.421095	-4.632179e+04
aco	68227	75005	5641.288782	-4.024887e+04
apalch	44634	70614	4579.546469	-6.258219e+04
oibdp	72863	75005	-3941.409192	-2.590904e+04
xint	64469	75005	-3829.352862	-8.317088e+03
caps	71931	75005	-3196.489599	-5.034086e+04
ivaeq	71457	75005	-3014.123461	-2.174282e+04
cstk	73080	75005	-3011.166145	-1.261688e+04
intan	74351	75005	-2919.453815	-1.103587e+04
recch	62416	72532	2604.026060	-2.102156e+04
ap	74917	75005	-2504.505248	-2.223392e+04
dltt	74817	75005	-2289.510691	-1.061091e+04
ppent	73282	75005	-2131.803875	-5.917365e+04
ceq	74856	75005	2099.646389	-2.799357e+03
dp	72283	75005	2094.861834	-1.117139e+04
invch	69690	73764	2027.820782	-2.817826e+04
oancf	74702	75005	-758.223668	-4.390364e+03
mibt	74163	75005	-579.344475	-3.365905e+03
ivncf	74689	75005	571.186625	-3.279110e+03
ibc	74551	75005	-487.204713	-2.187075e+03
csho	74201	75005	-333.528043	-2.694192e+04
dlcch	41862	70503	291.940128	-2.640138e+05
dlc	74974	75005	257.646866	-1.500792e+03
esubc	68207	75005	201.211062	-4.017797e+03
invt	74353	75005	177.616992	-4.859402e+03
chech	74701	74948	-148.009337	-3.349173e+03
tstk	74825	75005	-135.653129	-5.866449e+02
txach	52214	75005	-97.669890	-5.200987e+03
txdc	72533	75005	84.138237	-1.424961e+03
che	74999	75005	-70.437906	-6.643609e+02
fincf	74690	75005	-40.341627	-3.252400e+03
ao	75000	75005	-37.926984	-4.493027e+02
seq	75002	75005	-37.282301	-1.214421e+02
pstk	74786	75005	-20.953261	-3.301992e+02
pstkr	74930	75005	-7.086945	-8.506999e+01
nopi	74993	75005	6.593073	-6.715890e+01
txt	74990	75005	-5.734069	-2.316013e+01
pstkn	74794	75005	-4.220299	-2.150862e+02
aoloch	74543	74543	0.000000	0.000000e+00
dltr	73617	73617	0.000000	0.000000e+00
dltis	73249	73249	0.000000	0.000000e+00
cstke	75001	75001	0.000000	0.000000e+00

datadate	75005	75005	0.000000	0.000000e+00
at	75005	75005	0.000000	0.000000e+00
	delta_p50			
mkvalt	137.2548			
act	15675.0000			
lct	7207.0000			
lco	1865.5000			
rect	-189.0000			
re	-300.9000			
lt	46.0000			
aco	461.0000			
apalch	16.0000			
oibdp	182.6000			
xint	-68.5000			
caps	-443.0000			
ivaeq	0.0000			
cstk	5.0000			
intan	1.6000			
recch	9.0000			
ap	1.0000			
dltt	-21.0000			
ppent	519.5000			
ceq	27.5000			
dp	141.5300			
invch	0.0000			
oancf	-3.0000			
mibt	0.0000			
ivncf	6.0000			
ibc	-0.2400			
csho	-186.0000			
dlcch	0.0000			
dlc	3.0000			
esubc	0.0000			
invt	3.4000			
chech	-0.0900			
tstk	0.0000			
txach	0.0000			
txdc	0.0000			
che	-2.0000			
fincf	0.1000			
ao	-0.1000			
seq	-1.5000			
pstk	0.0000			
pstkr	0.0000			
nopi	0.0000			
txt	0.0000			
pstkn	0.0000			

aoloch	0.0000			
dltr	0.0000			
dltis	0.0000			
cstke	0.0000			
datadate	0.0000			
at	0.0000			
 == Change analysis: Imputed-only vs observed (pre) summary ==				
	n_imputed	imputed_mean	imputed_median	imputed_std \
col				
dlcch	33143	2.886764e+03	7.000000	6.776735e+04
apalch	30371	2.083248e+04	144.000000	7.569476e+04
txach	22791	0.000000e+00	0.000000	0.000000e+00
ivstch	19194	NaN	NaN	NaN
recch	12589	5.485655e+03	151.400000	8.040616e+04
mkvalt	12350	3.008425e+06	11457.197806	5.254540e+06
sppe	12239	NaN	NaN	NaN
act	10721	1.944047e+06	620683.448289	2.953969e+06
lct	10695	1.110617e+06	285127.913921	2.049787e+06
xint	10536	3.303843e+02	364.500000	1.277657e+02
txditc	9069	NaN	NaN	NaN
txp	7942	NaN	NaN	NaN
esubc	6798	0.000000e+00	0.000000	0.000000e+00
lco	6779	4.820648e+05	107133.519909	8.839821e+05
aco	6778	1.181587e+05	24795.542510	2.273255e+05
prcc_f	5956	NaN	NaN	NaN
prcc_c	5942	NaN	NaN	NaN
invch	5315	-1.884277e+03	0.000000	4.005203e+04
sppiv	3967	NaN	NaN	NaN
ivaeq	3548	5.011973e+03	0.000000	5.785298e+04
prstkc	3410	NaN	NaN	NaN
caps	3074	4.444179e+05	16696.000000	1.192783e+06
aqc	2812	NaN	NaN	NaN
dp	2722	1.376144e+05	53599.841951	2.450769e+05
re	2596	-7.659123e+04	-9263.000000	4.207461e+05
ivch	2540	NaN	NaN	NaN
ivao	2529	NaN	NaN	NaN
txdc	2472	0.000000e+00	0.000000	0.000000e+00
siv	2299	NaN	NaN	NaN
oibdp	2142	1.092507e+05	47840.710176	2.565068e+05
cstk	1925	2.676416e+04	428.596440	2.179249e+05
dpc	1757	NaN	NaN	NaN
dltis	1756	NaN	NaN	NaN
ppent	1723	4.942562e+05	237832.184614	7.872769e+05
sstk	1436	NaN	NaN	NaN
dltr	1388	NaN	NaN	NaN
dv	1137	NaN	NaN	NaN
spi	1126	NaN	NaN	NaN

rect	880	8.499478e+04	488.469851	3.713500e+05
mibt	842	3.299715e+01	0.000000	7.844159e+02
csho	804	1.943216e+05	11007.377133	4.821008e+05
dvt	706	NaN	NaN	NaN
dvc	703	NaN	NaN	NaN
intan	654	4.826358e+04	1666.201656	3.016515e+05
invt	652	1.675783e+05	1370.317109	5.706362e+05
txbcof	619	NaN	NaN	NaN
capx	503	NaN	NaN	NaN
xidoc	477	NaN	NaN	NaN
aoloch	462	NaN	NaN	NaN
ibc	454	-3.087908e+03	0.651928	3.557588e+04
		observed_mean_pre	observed_median_pre	\
col				
dlcch		2.168121e+03	0.0000	
apalch		8.385212e+03	50.0000	
txach		3.214308e+02	0.0000	
ivstch		6.871673e+04	0.0000	
recch		-1.318528e+04	-17.0000	
mkvalt		9.288907e+05	467.1767	
sppe		6.939724e+03	0.0000	
act		6.440956e+05	23074.0000	
lct		4.755088e+05	11405.0000	
xint		2.759126e+04	433.0000	
txditc		4.384545e+04	0.0000	
txp		1.062620e+04	0.0000	
esubc		-2.220041e+03	0.0000	
lco		2.039805e+05	3986.5000	
aco		5.573244e+04	946.0000	
prcc_f		6.838158e+02	13.9200	
prcc_c		7.067634e+02	14.0100	
invch		-3.860008e+04	0.0000	
sppiv		-7.943608e+03	0.0000	
ivaeq		6.873078e+04	0.0000	
prstkc		2.902337e+04	0.0000	
caps		5.224117e+05	26381.0000	
aqc		2.583727e+04	0.0000	
dp		7.989028e+04	1079.4700	
re		4.063079e+05	-1251.0000	
ivch		4.336588e+05	0.0000	
ivao		1.117387e+06	0.0000	
txdc		-2.552908e+03	0.0000	
siv		3.204806e+05	0.0000	
oibdp		2.472644e+05	800.4000	
cstk		1.440901e+05	197.0000	
dpc		8.009984e+04	1542.0000	
dltis		2.584506e+05	0.1500	

ppent	5.870571e+05	8837.5000
sstk	2.949253e+04	58.0000
dltr	2.679871e+05	140.7000
dv	4.155940e+04	0.0000
spi	-1.290533e+04	0.0000
rect	2.121967e+06	5841.0000
mibt	5.164076e+04	0.0000
csho	2.254364e+05	36145.0000
dvt	4.154428e+04	0.0000
dvc	4.081004e+04	0.0000
intan	3.830857e+05	1365.0000
invt	1.471455e+05	275.3000
txbcnf	2.081901e+02	0.0000
capx	8.862010e+04	887.0000
xidoc	8.325909e+02	0.0000
aoloch	3.689360e+04	-0.0100
ibc	7.740282e+04	2.0000

	mean_gap_imputed_minus_observed	median_gap_imputed_minus_observed
col		
dlcch	7.186430e+02	7.000000
apalch	1.244727e+04	94.000000
txach	-3.214308e+02	0.000000
ivstch	NaN	NaN
recch	1.867094e+04	168.400000
mkvalt	2.079534e+06	10990.021106
sppe	NaN	NaN
act	1.299952e+06	597609.448289
lct	6.351084e+05	273722.913921
xint	-2.726088e+04	-68.500000
txditc	NaN	NaN
txp	NaN	NaN
esubc	2.220041e+03	0.000000
lco	2.780843e+05	103147.019909
aco	6.242621e+04	23849.542510
prcc_f	NaN	NaN
prcc_c	NaN	NaN
invch	3.671580e+04	0.000000
sppiv	NaN	NaN
ivaeq	-6.371881e+04	0.000000
prstkc	NaN	NaN
caps	-7.799372e+04	-9685.000000
aqc	NaN	NaN
dp	5.772414e+04	52520.371951
re	-4.828992e+05	-8012.000000
ivch	NaN	NaN
ivao	NaN	NaN
txdc	2.552908e+03	0.000000

siv	NaN	NaN
oibdp	-1.380137e+05	47040.310176
cstk	-1.173260e+05	231.596440
dpc	NaN	NaN
dltis	NaN	NaN
ppent	-9.280090e+04	228994.684614
sstk	NaN	NaN
dltr	NaN	NaN
dv	NaN	NaN
spi	NaN	NaN
rect	-2.036972e+06	-5352.530149
mibt	-5.160776e+04	0.000000
csho	-3.111476e+04	-25137.622867
dvt	NaN	NaN
dvc	NaN	NaN
intan	-3.348221e+05	301.201656
invt	2.043276e+04	1095.017109
txbcof	NaN	NaN
capx	NaN	NaN
xidoc	NaN	NaN
aoloch	NaN	NaN
ibc	-8.049073e+04	-1.348072

```
[63]: # =====
# Feature Engineering & Target Construction
#   - Build leverage / coverage / cash-flow-to-debt ratios commonly used in ↴credit analysis
#   - Define a 'highly leveraged' distress proxy from multiple ratio-based ↴conditions
#   - Create the supervised-learning target: next-year distress within the same ↴firm
# =====
dlc = pd.to_numeric(df.get("dlc", np.nan), errors="coerce")
dltt = pd.to_numeric(df.get("dltt", np.nan), errors="coerce")
df["total_debt"] = pd.concat([dlc, dltt], axis=1).sum(axis=1, min_count=1)

seq = pd.to_numeric(df.get("seq", np.nan), errors="coerce")
mibt = pd.to_numeric(df.get("mibt", 0.0), errors="coerce")
df["equity_plus_mi_sp"] = seq + mibt
df["total_capital_sp"] = df["total_debt"] + df["equity_plus_mi_sp"]
df["sp_debt_to_capital"] = safe_divide(df["total_debt"], df["total_capital_sp"])

oibdp = pd.to_numeric(df.get("oibdp", np.nan), errors="coerce")
xint = pd.to_numeric(df.get("xint", np.nan), errors="coerce")
df["sp_debt_to_ebitda"] = safe_divide(df["total_debt"], oibdp)

txt = pd.to_numeric(df.get("txt", 0.0), errors="coerce").fillna(0.0)
```

```

txdc = pd.to_numeric(df.get("txdc", 0.0), errors="coerce").fillna(0.0)
txach = pd.to_numeric(df.get("txach", 0.0), errors="coerce").fillna(0.0)
df["cash_tax_paid_proxy"] = txt - txdc - txach

df["ffo_proxy"] = oibdp - xint - pd.to_numeric(df["cash_tax_paid_proxy"], □
    ↪errors="coerce")
df["sp_ffo_to_debt"] = safe_divide(df["ffo_proxy"], df["total_debt"])

oancf = pd.to_numeric(df.get("oancf", np.nan), errors="coerce")
capx = pd.to_numeric(df.get("capx", np.nan), errors="coerce")
df["sp_cfo_to_debt"] = safe_divide(oancf, df["total_debt"])
df["focf"] = oancf - capx
df["sp_focf_to_debt"] = safe_divide(df["focf"], df["total_debt"])

dv = pd.to_numeric(df.get("dv", 0.0), errors="coerce").fillna(0.0)
prstkc = pd.to_numeric(df.get("prstkc", 0.0), errors="coerce").fillna(0.0)
df["dcf"] = df["focf"] - dv - prstkc
df["sp_dcf_to_debt"] = safe_divide(df["dcf"], df["total_debt"])

# Log transforms (log1p handles 0 nicely). Negative → NaN.
for c in ["at", "mkvalt"]:
    if c in df.columns:
        s = pd.to_numeric(df[c], errors="coerce")
        df[f"log_{c}"] = np.where(s >= 0, np.log1p(s), np.nan)

# Interest coverage: EBITDA / |interest expense|
EPS_COV = 1e-6
df["sp_interest_coverage"] = safe_divide(oibdp, xint.abs() + EPS_COV)

# Distress proxy: 'highly leveraged' condition using multiple credit-ratio□
    ↪thresholds
td = pd.to_numeric(df["total_debt"], errors="coerce").to_numpy(dtype=float)
cap = pd.to_numeric(df["total_capital_sp"], errors="coerce").
    ↪to_numpy(dtype=float)
eb = pd.to_numeric(oibdp, errors="coerce").to_numpy(dtype=float)
ffo = pd.to_numeric(df["ffo_proxy"], errors="coerce").to_numpy(dtype=float)

ffo_to_debt_pct = 100.0 * safe_divide(ffo, td)
debt_to_capital_pct = 100.0 * safe_divide(td, cap)
debt_to_ebitda = safe_divide(td, eb)

# Three "highly leveraged" conditions (S&P table)
hl_ffo = (td > 0) & (ffo_to_debt_pct < 15.0)           # FFO/total debt < 15%
hl_cap = (cap > 0) & (debt_to_capital_pct > 55.0)       # Total debt/total□
    ↪capital > 55%
hl_deb = (td > 0) & (debt_to_ebitda > 4.5)      # TD/EBITDA > 4.5

```

```

is_highly_leveraged = hl_ffo & hl_cap & hl_deb

# Equity strictly negative and not missing
is_equity_negative = seq.notna() & (seq < -1)

# Final distress rule: highly leveraged (ratio-based proxy; excludes equity<0
# flag computed below)
df["distress_dummy"] = (is_highly_leveraged).astype("int8")

# Target: next year's distress (within firm)
df["target_next_year_distress"] = (
    df.groupby("firm_id")["distress_dummy"].shift(-1)
)
df = df.dropna(subset=["target_next_year_distress"]).reset_index(drop=True)
df["target_next_year_distress"] = df["target_next_year_distress"].astype("int8")

```

[64]:

```

# =====
# Train / Validation / Test Split (out-of-time)
#   - Split by label_year to respect the t → t+1 prediction structure
#   - Keep the last label year(s) inside the training pool as validation
# =====

train_pool = df[df["label_year"] <= TRAIN_CUTOFF_LABEL_YEAR].copy()
test = df[df["label_year"] > TRAIN_CUTOFF_LABEL_YEAR].copy()

years = np.sort(train_pool["label_year"].dropna().unique())
val_years = years[-VAL_YEARS:] if len(years) else np.array([], dtype=int)

val = train_pool[train_pool["label_year"].isin(val_years)].copy()
train = train_pool[~train_pool["label_year"].isin(val_years)].copy()

print(
    "Split:",
    f"train={len(train)}:{},",
    f"val={len(val)}:{},",
    f"test={len(test)}:{},",
    "| val_years:",
    list(val_years),
)

```

Split: train=44,783 val=6,415 test=12,404 | val_years: [np.int64(2022)]

[65]:

```

# =====
# Modeling-Ready Preprocessing (fit on TRAIN only)
#   - Handle infinities and remaining NaNs
#   - Winsorize features using TRAIN quantile bounds to reduce outlier leverage
#   - Standardize to z-scores using TRAIN mean/variance for comparable scaling
# =====

```

```

base_feats = [
    "sp_debt_to_capital",
    "sp_ffo_to_debt",
    "sp_cfo_to_debt",
    "sp_focf_to_debt",
    "sp_dcf_to_debt",
    "sp_debt_to_ebitda",
    "sp_interest_coverage",
    "log_at",
    "log_mkvalt",
]
feats = [c for c in base_feats if c in train.columns and c in val.columns and c in test.columns]

# Replace +/-inf with NaN
for d in (train, val, test):
    d[feats] = d[feats].replace([np.inf, -np.inf], np.nan)

# Remaining NaNs: median imputation fit on TRAIN only (after ratio construction)
fill = train[feats].median(numeric_only=True)
for d in (train, val, test):
    d[feats] = d[feats].fillna(fill)

# Winsorize each feature using train-only quantiles
bounds = {}
for c in feats:
    s = pd.to_numeric(train[c], errors="coerce")
    bounds[c] = (s.quantile(WINSOR_LOWER_Q), s.quantile(WINSOR_UPPER_Q))

for d in (train, val, test):
    for c, (lo, hi) in bounds.items():
        s = pd.to_numeric(d[c], errors="coerce")
        d[c] = s.clip(lo, hi)

# Standardize (z-scores), fit on train only
x_train = train[feats].to_numpy(dtype=float)
x_val = val[feats].to_numpy(dtype=float)
x_test = test[feats].to_numpy(dtype=float)

scaler = StandardScaler().fit(x_train)
x_train_z = scaler.transform(x_train)
x_val_z = scaler.transform(x_val)
x_test_z = scaler.transform(x_test)

z_cols = [f"z_{c}" for c in feats]
train[z_cols] = x_train_z
val[z_cols] = x_val_z

```

```

test[z_cols] = x_test_z

[66]: # =====
# Diagnostics & Monitoring Proxies
#   - Correlation screen (TRAIN) for rough signal strength and sanity checks
#   - Expanding-window time folds for temporal stability checks
#   - Dataset overview (rows/firms/years/target rate) + target rate by year
#   - Distribution summaries, collinearity scan, and simple drift proxy (SMD)
#     ↳ Train→Test
# =====
t = "target_next_year_distress"

corr = (
    train[[t] + feats]
    .corr(numeric_only=True)[t]
    .drop(t)
    .sort_values(key=np.abs, ascending=False)
)
print("Correlation with target:")
print(corr)

# Multicollinearity: Variance Inflation Factor (VIF)
def calculate_vif(df, features):
    vif_data = pd.DataFrame()
    vif_data["feature"] = features
    vif_data["VIF"] = [variance_inflation_factor(df[features].values, i) for i in range(len(features))]
    return vif_data.sort_values("VIF", ascending=False)

vif_df = calculate_vif(train, z_cols)
print("\n==== Multicollinearity Diagnostic (VIF) ===")
print(vif_df)

folds = rolling_year_folds(train_pool, n_splits=N_SPLITS_TIME_CV, min_train_years=3)
for i, (tr_idx, va_idx, tr_years, va_year) in enumerate(folds, 1):
    print(
        f"Fold {i}: train_years={tr_years[0]}..{tr_years[-1]}, "
        f"n={len(tr_idx)}, "
        f"val_year={va_year} (n={len(va_idx)})"
    )

def _overview(d: pd.DataFrame, name: str) -> None:
    n_rows = len(d)
    n_firms = d["firm_id"].nunique() if "firm_id" in d.columns else np.nan
    n_years = d["fyear"].nunique() if "fyear" in d.columns else np.nan

```

```

target_rate = float(d[t].mean()) if t in d.columns else np.nan

print(f"\n== {name} == rows={n_rows:,} | firms={n_firms:,} | "
      f"years={n_years} | target_rate={target_rate:.4f}")

if "label_year" in d.columns:
    by_year = d.groupby("label_year")[t].agg(["mean", "count"])
    print("\nTarget by label_year (tail):")
    print(by_year.tail(12))

_overview(train, "TRAIN")
_overview(val, "VAL")
_overview(test, "TEST")

post_miss = pd.DataFrame({"col": raw})
post_miss["train_pct_na"] = [train[c].isna().mean() * 100 if c in train.columns
                             else np.nan for c in raw]
post_miss["val_pct_na"] = [val[c].isna().mean() * 100 if c in val.columns
                           else np.nan for c in raw]
post_miss["test_pct_na"] = [test[c].isna().mean() * 100 if c in test.columns
                            else np.nan for c in raw]
if not post_miss.empty:
    post_miss = post_miss.sort_values("train_pct_na", ascending=False)
    print("\nPost-imputation missingness on raw inputs (pct):")
    print(post_miss.head(50).round(4))

def _dist(d: pd.DataFrame, cols: list[str], name: str) -> pd.DataFrame:
    x = d[cols].replace([np.inf, -np.inf], np.nan)
    q = x.quantile([0.01, 0.05, 0.25, 0.50, 0.75, 0.95, 0.99]).T
    out = pd.DataFrame(
        {
            "n": x.notna().sum(),
            "mean": x.mean(),
            "std": x.std(ddof=0),
            "min": x.min(),
            "p01": q[0.01],
            "p05": q[0.05],
            "p25": q[0.25],
            "p50": q[0.50],
            "p75": q[0.75],
            "p95": q[0.95],
            "p99": q[0.99],
            "max": x.max(),
            "skew": x.skew(numeric_only=True),
            "kurt": x.kurtosis(numeric_only=True),
        }
    )

```

```

    )
    print(f"\nDistribution summary ({name})")
    print(out.round(4).sort_values("skew", key=lambda s: s.abs(), u
↪ascending=False))
    return out

_ = _dist(train, feats, "TRAIN | winsorized raw feats")
_ = _dist(train, z_cols, "TRAIN | standardized feats")

def _hi_corr(d: pd.DataFrame, cols: list[str], thr: float = 0.80) -> u
↪list[tuple[str, str, float]]:
    cm = d[cols].corr(numeric_only=True)
    pairs = []
    for i in range(len(cols)):
        for j in range(i + 1, len(cols)):
            r = cm.iloc[i, j]
            if np.isfinite(r) and abs(r) >= thr:
                pairs.append((cols[i], cols[j], float(r)))
    return sorted(pairs, key=lambda x: abs(x[2]), reverse=True)

pairs = _hi_corr(train, feats, thr=0.80)
print("\nHigh collinearity pairs among feats (|corr|>=0.80) [top 25]:")
for a, b, r in pairs[:25]:
    print(f"{a} vs {b}: r={r:.3f}")

def _drift_smd(a_df: pd.DataFrame, b_df: pd.DataFrame, cols: list[str]) -> pd.
↪DataFrame:
    rows = []
    for c in cols:
        a = pd.to_numeric(a_df[c], errors="coerce").replace([np.inf, -np.inf], u
↪np.nan)
        b = pd.to_numeric(b_df[c], errors="coerce").replace([np.inf, -np.inf], u
↪np.nan)

        ma, mb = float(a.mean()), float(b.mean())
        sa, sb = float(a.std(ddof=0)), float(b.std(ddof=0))
        sp = np.sqrt(0.5 * (sa ** 2 + sb ** 2))
        smd = (mb - ma) / sp if sp > 0 else np.nan

        rows.append((c, ma, mb, smd, abs(smd) if np.isfinite(smd) else np.nan))

    out = pd.DataFrame(rows, columns=["feature", "mean_train", "mean_test", u
↪"smd", "abs_smd"])

```

```

    return out.sort_values("abs_smd", ascending=False)

drift = _drift_smd(train, test, feats)
print("\nTrain→Test drift (SMD) [top 15]:")
print(drift.head(15).round(4))

def _group_diff(d: pd.DataFrame, cols: list[str]) -> pd.Series:
    g = d.groupby(t)[cols].mean(numeric_only=True)
    if 0 in g.index and 1 in g.index:
        return (g.loc[1] - g.loc[0]).sort_values(key=np.abs, ascending=False)
    return pd.Series(dtype="float64")

diff = _group_diff(train, feats)
if not diff.empty:
    print("\nMean difference (target=1 minus target=0) on TRAIN feats [top 15]:")
    print(diff.head(15).round(4))

```

Correlation with target:

sp_debt_to_capital	0.239281
log_at	0.148383
sp_debt_to_ebitda	0.097794
log_mkvalt	0.094928
sp_dcf_to_debt	0.039173
sp_focf_to_debt	0.025729
sp_ffo_to_debt	0.021727
sp_interest_coverage	0.006064
sp_cfo_to_debt	-0.001887

Name: target_next_year_distress, dtype: float64

==== Multicollinearity Diagnostic (VIF) ===

	feature	VIF
3	z_sp_focf_to_debt	10.773549
4	z_sp_dcf_to_debt	6.370928
2	z_sp_cfo_to_debt	3.514626
1	z_sp_ffo_to_debt	2.131804
7	z_log_at	1.281938
8	z_log_mkvalt	1.243078
0	z_sp_debt_to_capital	1.060240
5	z_sp_debt_to_ebitda	1.012524
6	z_sp_interest_coverage	1.005194

Fold 1: train_years=2015..2017 (n=19775), val_year=2018 (n=6337)
 Fold 2: train_years=2015..2018 (n=26112), val_year=2019 (n=6173)
 Fold 3: train_years=2015..2019 (n=32285), val_year=2020 (n=6233)
 Fold 4: train_years=2015..2020 (n=38518), val_year=2021 (n=6265)

```
==== TRAIN === rows=44,783 | firms=9,220 | years=7 | target_rate=0.1076
```

```
Target by label_year (tail):
```

	mean	count
label_year		
2015	0.104385	6773
2016	0.104566	6570
2017	0.100280	6432
2018	0.101310	6337
2019	0.124737	6173
2020	0.123376	6233
2021	0.095611	6265

```
==== VAL === rows=6,415 | firms=6,415 | years=1 | target_rate=0.0929
```

```
Target by label_year (tail):
```

	mean	count
label_year		
2022	0.092907	6415

```
==== TEST === rows=12,404 | firms=6,633 | years=2 | target_rate=0.1051
```

```
Target by label_year (tail):
```

	mean	count
label_year		
2023	0.105579	6327
2024	0.104657	6077

```
Post-imputation missingness on raw inputs (pct):
```

	col	train_pct_na	val_pct_na	test_pct_na
48	ivstch	23.4218	30.1949	30.3773
71	sppe	16.6559	16.3055	16.1641
78	txditc	12.4668	11.9719	12.1251
79	txp	10.9283	10.4910	10.7546
60	prcc_f	8.8516	6.1574	7.2718
59	prcc_c	8.8493	6.1107	7.2557
18	dlcch	8.2219	3.5230	2.0558
5	apalch	7.8802	3.1489	2.4105
61	prstkc	4.9997	4.4895	3.3538
72	sppiv	4.9997	5.5495	5.9174
66	recch	4.0283	2.2603	2.0880
46	ivch	3.4723	3.3359	3.5069
6	aqc	3.3227	4.6921	4.5953
45	ivao	3.2646	3.4451	3.7730
69	siv	3.1731	2.9774	2.9507
19	dltis	2.4585	2.3071	1.9832
24	dpc	2.4250	2.2136	2.1928

20	dltr	2.0276	1.2627	1.6608
40	invch	1.9985	1.0444	0.9755
73	sstk	1.8869	1.9174	2.0477
25	dv	1.5899	1.4030	1.3463
70	spi	1.4112	1.6680	1.8381
76	txbcof	1.0004	0.4677	0.4595
28	dvt	0.8374	0.5924	0.4837
26	dvc	0.8329	0.5924	0.4837
9	capx	0.7079	0.4365	0.5160
83	xidoc	0.6967	0.4677	0.4434
3	aoloch	0.6677	0.4677	0.4353
33	fopo	0.6319	0.4209	0.3950
30	exre	0.5940	0.3118	0.2902
31	fiao	0.4868	0.3118	0.2660
43	ivaco	0.4712	0.2962	0.2660
12	chech	0.1072	0.0156	0.0000
27	dvp	0.0536	0.0468	0.0484
15	cstke	0.0089	0.0000	0.0000
82	xido	0.0067	0.0000	0.0000
81	xi	0.0022	0.0000	0.0000
22	do	0.0022	0.0000	0.0000
8	caps	0.0000	0.0000	0.0000
7	at	0.0000	0.0000	0.0000
1	act	0.0000	0.0000	0.0000
23	dp	0.0000	0.0000	0.0000
21	dltt	0.0000	0.0000	0.0000
17	dlc	0.0000	0.0000	0.0000
10	ceq	0.0000	0.0000	0.0000
11	che	0.0000	0.0000	0.0000
13	csho	0.0000	0.0000	0.0000
14	cstk	0.0000	0.0000	0.0000
16	datadate	0.0000	0.0000	0.0000
4	ap	0.0000	0.0000	0.0000

Distribution summary (TRAIN | winsorized raw feats)

	n	mean	std	min	\		
sp_dcf_to_debt	44783	-2.502190e+01	1.984466e+02	-1.653794e+03			
sp_debt_to_ebitda	44783	1.314641e+02	8.146559e+02	-8.128037e+02			
sp_focf_to_debt	44783	-1.282550e+01	1.741842e+02	-1.360862e+03			
sp_ffo_to_debt	44783	-1.100000e+01	2.290823e+02	-1.794695e+03			
sp_interest_coverage	44783	-6.098563e+07	6.492082e+09	-3.280108e+10			
sp_debt_to_capital	44783	3.126000e-01	5.945000e-01	-2.786200e+00			
log_mkvalt	44783	6.999300e+00	4.093800e+00	3.656000e-01			
log_at	44783	1.107300e+01	3.565000e+00	6.931000e-01			
sp_cfo_to_debt	44783	4.686000e+00	1.540910e+02	-9.350585e+02			
		p01	p05	p25	p50	p75	\
sp_dcf_to_debt		-1.650683e+03	-3.161480e+01	-0.2674	-0.0163	0.0919	

sp_debt_to_ebitda	-8.108731e+02	-6.313900e+00	0.0000	0.3204	3.6265
sp_focf_to_debt	-1.360587e+03	-2.001020e+01	-0.1546	0.0470	0.1921
sp_ffo_to_debt	-1.793826e+03	-1.929870e+01	-0.0832	0.1162	0.3252
sp_interest_coverage	-3.278543e+10	-2.221000e+09	-7.5590	2.7042	21.0105
sp_debt_to_capital	-2.782800e+00	-1.000000e-04	0.0002	0.2456	0.5571
log_mkvalt	3.656000e-01	1.299500e+00	3.9637	6.3919	9.1257
log_at	6.931000e-01	4.369400e+00	8.8541	11.4487	13.7914
sp_cfo_to_debt	-9.345454e+02	-7.473200e+00	0.0001	0.1332	0.3434

	p95	p99	max	skew	kurt
sp_dcf_to_debt	3.0767	3.510140e+02	3.512747e+02	-6.8466	51.0331
sp_debt_to_ebitda	39.9676	6.328203e+03	6.332345e+03	6.3041	41.3162
sp_focf_to_debt	6.6901	6.072581e+02	6.073161e+02	-5.1643	41.3521
sp_ffo_to_debt	14.5162	8.582689e+02	8.585733e+02	-4.8162	41.3640
sp_interest_coverage	37063.0105	4.293433e+10	4.300872e+10	1.7497	27.9326
sp_debt_to_capital	1.0013	2.785900e+00	2.787600e+00	-0.6996	10.4704
log_mkvalt	15.6653	1.645580e+01	1.645580e+01	0.6659	-0.2189
log_at	16.0542	1.785350e+01	1.785370e+01	-0.6152	0.1416
sp_cfo_to_debt	17.0351	9.284942e+02	9.291911e+02	0.2574	28.3851

Distribution summary (TRAIN | standardized feats)

	n	mean	std	min	p01	p05	p25	\
z_sp_dcf_to_debt	44783	-0.0	1.0	-8.2076	-8.1919	-0.0332	0.1247	
z_sp_debt_to_ebitda	44783	-0.0	1.0	-1.1591	-1.1567	-0.1691	-0.1614	
z_sp_focf_to_debt	44783	-0.0	1.0	-7.7391	-7.7376	-0.0412	0.0727	
z_sp_ffo_to_debt	44783	0.0	1.0	-7.7863	-7.7825	-0.0362	0.0477	
z_sp_interest_coverage	44783	-0.0	1.0	-5.0431	-5.0407	-0.3327	0.0094	
z_sp_debt_to_capital	44783	-0.0	1.0	-5.2125	-5.2068	-0.5260	-0.5255	
z_log_mkvalt	44783	-0.0	1.0	-1.6204	-1.6204	-1.3923	-0.7415	
z_log_at	44783	0.0	1.0	-2.9116	-2.9116	-1.8804	-0.6224	
z_sp_cfo_to_debt	44783	-0.0	1.0	-6.0986	-6.0953	-0.0789	-0.0304	

	p50	p75	p95	p99	max	skew	\
z_sp_dcf_to_debt	0.1260	0.1266	0.1416	1.8949	1.8962	-6.8466	
z_sp_debt_to_ebitda	-0.1610	-0.1569	-0.1123	7.6066	7.6117	6.3041	
z_sp_focf_to_debt	0.0739	0.0747	0.1120	3.5599	3.5603	-5.1643	
z_sp_ffo_to_debt	0.0485	0.0494	0.1114	3.7946	3.7959	-4.8162	
z_sp_interest_coverage	0.0094	0.0094	0.0094	6.6227	6.6342	1.7497	
z_sp_debt_to_capital	-0.1127	0.4113	1.1584	4.1602	4.1631	-0.6996	
z_log_mkvalt	-0.1484	0.5194	2.1168	2.3100	2.3100	0.6659	
z_log_at	0.1054	0.7625	1.3973	1.9020	1.9020	-0.6152	
z_sp_cfo_to_debt	-0.0295	-0.0282	0.0801	5.9952	5.9997	0.2574	

	kurt
z_sp_dcf_to_debt	51.0331
z_sp_debt_to_ebitda	41.3162
z_sp_focf_to_debt	41.3521
z_sp_ffo_to_debt	41.3640

```

z_sp_interest_coverage 27.9326
z_sp_debt_to_capital    10.4704
z_log_mkvalt           -0.2189
z_log_at                0.1416
z_sp_cfo_to_debt       28.3851

High collinearity pairs among feats (|corr|>=0.80) [top 25]:
sp_focf_to_debt vs sp_dcf_to_debt: r=0.897

```

Train→Test drift (SMD) [top 15]:

	feature	mean_train	mean_test	smd	abs_smd
6	sp_interest_coverage	-6.098563e+07	-7.379267e+08	-0.1040	0.1040
2	sp_cfo_to_debt	4.686000e+00	-7.884400e+00	-0.0902	0.0902
7	log_at	1.107300e+01	1.136440e+01	0.0836	0.0836
0	sp_debt_to_capital	3.126000e-01	3.574000e-01	0.0771	0.0771
8	log_mkvalt	6.999300e+00	7.247000e+00	0.0606	0.0606
1	sp_ffo_to_debt	-1.100000e+01	-1.968880e+01	-0.0407	0.0407
3	sp_focf_to_debt	-1.282550e+01	-1.753580e+01	-0.0287	0.0287
4	sp_dcf_to_debt	-2.502190e+01	-2.412000e+01	0.0048	0.0048
5	sp_debt_to_ebitda	1.314641e+02	1.284517e+02	-0.0037	0.0037

Mean difference (target=1 minus target=0) on TRAIN feats [top 15]:

```

sp_interest_coverage     1.270463e+08
sp_debt_to_ebitda      2.570901e+02
sp_dcf_to_debt         2.508610e+01
sp_ffo_to_debt         1.606200e+01
sp_focf_to_debt        1.446220e+01
log_at                 1.707100e+00
log_mkvalt              1.254100e+00
sp_cfo_to_debt         -9.383000e-01
sp_debt_to_capital     4.591000e-01
dtype: float64

```

```
[67]: # =====
# Visual EDA: Feature distributions by distress flag
#   - Quick separation check: do distressed vs non-distressed firms differ in
#     ↪ levels?
#   - Uses a horizontal boxplot for comparability across features
# =====

import matplotlib.pyplot as plt
import seaborn as sns

# Objective: visualize feature distribution differences for distressed vs
# ↪ non-distressed observations.

plot_df = train_pool.copy() if "train_pool" in globals() else df.copy()
```

```

flag_col = "distress_dummy" if "distress_dummy" in plot_df.columns else (
    "target_next_year_distress" if "target_next_year_distress" in plot_df.
    ↪columns else None
)
if flag_col is None:
    raise KeyError("No distress flag found. Expected 'distress_dummy' or
    ↪'target_next_year_distress' in the data.")

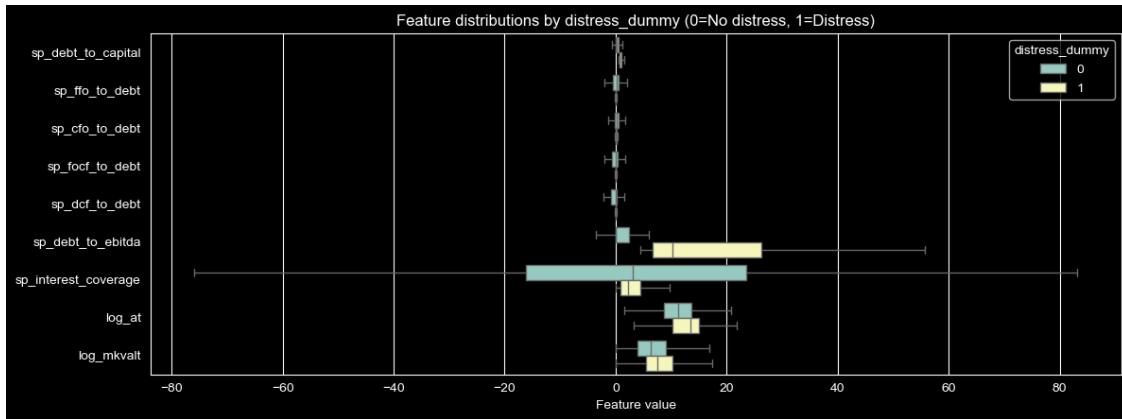
plot_feats = [c for c in (feats if "feats" in globals() else []) if c in
    ↪plot_df.columns]
if not plot_feats:
    plot_feats = [c for c in (z_cols if "z_cols" in globals() else []) if c in
    ↪plot_df.columns]
if not plot_feats:
    raise KeyError(
        "No feature columns found to plot. Expected 'feats' or 'z_cols' to
        ↪exist and be present in the data.")

tmp = plot_df[[flag_col] + plot_feats].copy()
tmp[flag_col] = pd.to_numeric(tmp[flag_col], errors="coerce").astype("Int64")
tmp = tmp[tmp[flag_col].isin([0, 1])].copy()

long = tmp.melt(id_vars=[flag_col], value_vars=plot_feats, var_name="feature",
    ↪value_name="value")
long["value"] = pd.to_numeric(long["value"], errors="coerce")
long = long.dropna(subset=["value"])

fig, ax = plt.subplots(figsize=(12, max(4.5, 0.45 * len(plot_feats))))
sns.boxplot(
    data=long,
    x="value",
    y="feature",
    hue=flag_col,
    orient="h",
    showfliers=False,
    ax=ax,
)
ax.set_title(f"Feature distributions by {flag_col} (0=No distress, 1=Distress)")
ax.set_xlabel("Feature value")
ax.set_ylabel("")
ax.legend(title=flag_col, loc="best")
plt.tight_layout()
plt.show()

```



```
[68]: # =====
# Sanity checks on the distress proxy
#   - Component prevalence (each condition and joint condition)
#   - Distress rate by firm size decile (log assets) to check monotonic patterns
# =====

print(pd.Series({
    "hl_ffo": hl_ffo.mean(),
    "hl_cap": hl_cap.mean(),
    "hl_deb": hl_deb.mean(),
    "hl_all": is_highly_leveraged.mean()
}))
df["size_decile"] = pd.qcut(df["log_at"], 10, duplicates="drop")
print(df.groupby("size_decile")["distress_dummy"].mean())
```

```
hl_ffo      0.481235
hl_cap     0.264142
hl_deb     0.204013
hl_all     0.103446
dtype: float64
size_decile
(-0.001, 6.585]      0.029241
(6.585, 8.38]        0.093553
(8.38, 9.509]        0.094782
(9.509, 10.52]       0.057408
(10.52, 11.507]      0.053459
(11.507, 12.554]     0.073585
(12.554, 13.434]     0.098270
(13.434, 14.245]     0.143396
(14.245, 15.252]     0.182704
(15.252, 21.886]     0.208458
Name: distress_dummy, dtype: float64
```

```
/var/folders/p1/_cwwbdbj51q1lwpynfnzdpxm0000gn/T/ipykernel_2951/2342790164.py:14
: FutureWarning: The default of observed=False is deprecated and will be changed
to True in a future version of pandas. Pass observed=False to retain current
behavior or observed=True to adopt the future default and silence this warning.
    print(df.groupby("size_decile")["distress_dummy"].mean())
```

0.0.1 5.5 Endogeneity & Simultaneity Bias: A Critical Note

As highlighted by recent literature (e.g., Roberts 2015, Jiang 2017), modeling capital structure as a predictor of credit outcomes is prone to **endogeneity**.

1. **Simultaneity:** Firms choose their leverage (Debt/Capital, etc.) based on their expected future credit risk and investment opportunities. Thus, leverage and distress outcomes are determined simultaneously in equilibrium.
2. **Reverse Causality:** While high leverage may lead to distress, the *anticipation* of distress may also force firms to seek more debt (if they follow pecking order) or prevent them from accessing debt (if credit constrained).

Mitigation in this notebook: - While we use a standard Logistic Regression for transparency and baseline benchmarking, we acknowledge that the coefficients should be interpreted as **predictive associations** rather than pure causal effects. - Future iterations should explore **Instrumental Variables (IV)** or **GMM** estimation (e.g., using industry-median leverage or historical tax changes as instruments) to better identify the causal impact of leverage on distress.

```
[69]: # =====
# Persistence Benchmark (Economic Sanity Check)
#   - Does the ML model add value over simply assuming distress persists?
#   - We use 'distress_dummy' (current year) to predict
#     ↪ 'target_next_year_distress'
# =====
from sklearn.metrics import f1_score, accuracy_score

for name, d in [("VAL", val), ("TEST", test)]:
    y_true = d[TARGET_COL].astype(int)
    y_bench = d["distress_dummy"].astype(int)

    auc_bench = roc_auc_score(y_true, y_bench)
    f1_bench = f1_score(y_true, y_bench)
    acc_bench = accuracy_score(y_true, y_bench)

    print(f"--- Persistence Benchmark ({name}) ---")
    print(f"AUC: {auc_bench:.4f} | F1: {f1_bench:.4f} | Accuracy: {acc_bench:.
         .4f}")
    print("")

--- Persistence Benchmark (VAL) ---
AUC: 0.7680 | F1: 0.5798 | Accuracy: 0.9221
--- Persistence Benchmark (TEST) ---
AUC: 0.7715 | F1: 0.6036 | Accuracy: 0.9193
```

0.1 6. Modeling: Baseline out-of-sample supervised learning

This section trains a simple, interpretable classifier on the **TRAIN** split, selects a small hyperparameter setting using the **VAL** split, and reports final performance on the **TEST** split (held out, `label_year > train cutoff`).

Goal: Predict `target_next_year_distress` one year ahead from standardized financial-ratio features.

```
[70]: # =====
# 6.1 Setup: Features, target, and train/val/test matrices
# =====
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    roc_auc_score,
    average_precision_score,
    brier_score_loss,
    confusion_matrix,
    classification_report,
    precision_recall_curve,
    roc_curve,
)
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from IPython.display import display

TARGET_COL = "target_next_year_distress"

# Prefer standardized feature columns created earlier (fit on TRAIN only)
if "z_cols" in globals():
    MODEL_FEATS = list(z_cols)
else:
    MODEL_FEATS = [f"z_{c}" for c in feats] # fallback

# Basic sanity checks
for df_name, df_ in [("train", train), ("val", val), ("test", test)]:
    missing_feats = [c for c in MODEL_FEATS if c not in df_.columns]
    if missing_feats:
        raise KeyError(f"{df_name}: missing feature columns: {missing_feats}")

X_train = train[MODEL_FEATS].to_numpy(dtype=float)
y_train = train[TARGET_COL].astype(int).to_numpy()

X_val = val[MODEL_FEATS].to_numpy(dtype=float)
y_val = val[TARGET_COL].astype(int).to_numpy()

X_test = test[MODEL_FEATS].to_numpy(dtype=float)
```

```

y_test = test[TARGET_COL].astype(int).to_numpy()

# Defensive: ensure model inputs are finite
def _assert_finite(name, X):
    bad = ~np.isfinite(X)
    if bad.any():
        rows, cols = np.where(bad)
        raise ValueError(f"{name}: found non-finite values at {len(rows)} cells"
                         f"(e.g., row={rows[0]}, col={MODEL_FEATS[cols[0]]}).")

_assert_finite("X_train", X_train)
_assert_finite("X_val", X_val)
_assert_finite("X_test", X_test)

print("Modeling matrix shapes:")
print(f" X_train: {X_train.shape} | y_train mean={y_train.mean():.4f}")
print(f" X_val:   {X_val.shape} | y_val mean={y_val.mean():.4f}")
print(f" X_test:  {X_test.shape} | y_test mean={y_test.mean():.4f}")

```

Modeling matrix shapes:

```

X_train: (44783, 9) | y_train mean=0.1076
X_val:   (6415, 9) | y_val mean=0.0929
X_test:  (12404, 9) | y_test mean=0.1051

```

```

[71]: # =====
# 6.2 Baseline model: Logistic Regression (tuned on VAL)
#   - Interpretable, strong baseline for tabular finance ratios
#   - We tune C on VAL; you can expand the grid later if needed
# =====
C_GRID = [0.01, 0.1, 1.0, 10.0]

best = {"C": None, "val_auc": -np.inf, "model": None}

for C in C_GRID:
    clf = LogisticRegression(
        C=C,
        solver="lbfgs",
        max_iter=500,
        class_weight="balanced", # practical default with ~10% positives
        n_jobs=None,
        random_state=42,
    )
    clf.fit(X_train, y_train)
    val_proba = clf.predict_proba(X_val)[:, 1]
    val_auc = roc_auc_score(y_val, val_proba)

    if val_auc > best["val_auc"]:

```

```

        best.update({"C": C, "val_auc": val_auc, "model": clf})

print(f"Best LogisticRegression on VAL: C={best['C']} | AUC={best['val_auc']:.4f}")
logit = best["model"]

```

Best LogisticRegression on VAL: C=10.0 | AUC=0.8022

```
[72]: # =====
# 6.2b Statistical Inference & Economic Interpretation (Statsmodels)
# - sklearn is great for prediction but lacks p-values and inference.
# - We re-estimate the logit using statsmodels to audit statistical significance.
# - We also report Marginal Effects to see economic significance.
# =====

import statsmodels.api as sm

# Add constant for intercept
X_train_sm = sm.add_constant(X_train)
sm_model = sm.Logit(y_train, X_train_sm).fit(disp=0)

print("\n==== Statsmodels Logistic Regression Summary ===")
print(sm_model.summary(xname=["const"] + MODEL_FEATS))

# Marginal Effects at the Mean (MEM)
try:
    mfx = sm_model.get_margeff(at="mean")
    print("\n==== Economic Significance: Marginal Effects at the Mean ===")
    print(mfx.summary())
except Exception as e:
    print(f"\nCould not calculate marginal effects: {e}")

# AIC / BIC / Pseudo-R2
print("\nModel Fit:")
print(f"  Pseudo R2 (McFadden): {sm_model.prsquared:.4f}")
print(f"  AIC: {sm_model.aic:.2f}")
print(f"  BIC: {sm_model.bic:.2f}")

```

```

==== Statsmodels Logistic Regression Summary ===
          Logit Regression Results
=====
Dep. Variable:                  y      No. Observations:      44783
Model:                          Logit   Df Residuals:           44773
Method:                         MLE    Df Model:                 9
Date:  Sat, 27 Dec 2025       Pseudo R-squ.:      0.1287
Time: 21:58:01                Log-Likelihood: -13324.
converged:                      True   LL-Null:      -15293.

```

Covariance Type:	nonrobust	LLR p-value:	0.000			
<hr/>						
<hr/>						
	coef	std err	z	P> z	[0.025	
0.975]						
<hr/>						
const	-2.487	-2.5290	0.021	-118.866	0.000	-2.571
z_sp_debt_to_capital	0.868	0.8344	0.017	48.294	0.000	0.800
z_sp_ffo_to_debt	0.238	0.1321	0.054	2.447	0.014	0.026
z_sp_cfo_to_debt	-0.082	-0.2272	0.074	-3.060	0.002	-0.373
z_sp_focf_to_debt	0.180	-0.1136	0.150	-0.758	0.448	-0.408
z_sp_dcf_to_debt	0.593	0.3571	0.120	2.969	0.003	0.121
z_sp_debt_to_ebitda	0.159	0.1376	0.011	12.428	0.000	0.116
z_sp_interest_coverage	0.033	-0.0062	0.020	-0.312	0.755	-0.045
z_log_at	0.565	0.5240	0.021	25.229	0.000	0.483
z_log_mkvalt	0.135	0.1008	0.018	5.750	0.000	0.066
<hr/>						
<hr/>						
==== Economic Significance: Marginal Effects at the Mean ===						
Logit Marginal Effects						
<hr/>						
Dep. Variable:	y					
Method:	dydx					
At:	mean					
<hr/>						
	dy/dx	std err	z	P> z	[0.025	0.975]
x1	0.0571	0.001	47.717	0.000	0.055	0.059
x2	0.0090	0.004	2.462	0.014	0.002	0.016
x3	-0.0155	0.005	-3.077	0.002	-0.025	-0.006
x4	-0.0078	0.010	-0.758	0.448	-0.028	0.012
x5	0.0244	0.008	2.989	0.003	0.008	0.040
x6	0.0094	0.001	12.295	0.000	0.008	0.011
x7	-0.0004	0.001	-0.312	0.755	-0.003	0.002
x8	0.0358	0.001	26.763	0.000	0.033	0.038
x9	0.0069	0.001	5.762	0.000	0.005	0.009

=====
Model Fit:

Pseudo R2 (McFadden): 0.1287
AIC: 26668.28
BIC: 26755.37

```
[76]: # =====  
# 6.2c Temporal Stability (Walk-Forward Validation)  
#   - We assess whether model performance is stable across different time  
#     regimes.  
#   - Using the expanding-window folds defined in diagnostics.  
# =====  
  
# Ensure MODEL_FEATS exist in train_pool (train_pool may not have z_* columns)  
_missing = [c for c in MODEL_FEATS if c not in train_pool.columns]  
if _missing:  
    # Recreate the exact same preprocessing pipeline used for train/val/test,  
    # but applied to train_pool  
    for c in feats:  
        train_pool[c] = pd.to_numeric(train_pool[c], errors="coerce").  
        replace([np.inf, -np.inf], np.nan)  
        train_pool[feats] = train_pool[feats].fillna(fill)  
        for c, (lo, hi) in bounds.items():  
            train_pool[c] = pd.to_numeric(train_pool[c], errors="coerce").clip(lo,  
            hi)  
  
        x_pool = train_pool[feats].to_numpy(dtype=float)  
        x_pool_z = scaler.transform(x_pool)  
        train_pool[z_cols] = x_pool_z  
  
temporal_results = []  
for i, (tr_idx, va_idx, tr_years, va_year) in enumerate(folds, 1):  
    # Prepare fold data (fold indices are label-based, so use .loc not .iloc)  
    X_tr = train_pool.loc[tr_idx, MODEL_FEATS].to_numpy(dtype=float)  
    y_tr = train_pool.loc[tr_idx, TARGET_COL].astype(int).to_numpy()  
    X_va = train_pool.loc[va_idx, MODEL_FEATS].to_numpy(dtype=float)  
    y_va = train_pool.loc[va_idx, TARGET_COL].astype(int).to_numpy()  
  
    # Fit model (using best C from tuning)  
    fold_clf = LogisticRegression(C=best["C"], solver="lbfgs", max_iter=500,  
    class_weight="balanced", random_state=42)  
    fold_clf.fit(X_tr, y_tr)  
  
    # Evaluate  
    probs = fold_clf.predict_proba(X_va)[:, 1]  
    auc = roc_auc_score(y_va, probs)
```

```

ap = average_precision_score(y_va, probs)

temporal_results.append({
    "Fold": i,
    "Val_Year": va_year,
    "Train_End": tr_years[-1],
    "AUC": auc,
    "PR-AUC": ap
})

temporal_df = pd.DataFrame(temporal_results)
print("\n==== Temporal Stability: Walk-Forward Results ===")
print(temporal_df.round(4))
if not temporal_df.empty:
    print(f"\nAverage AUC across folds: {temporal_df['AUC'].mean():.4f}")

```

```

==== Temporal Stability: Walk-Forward Results ===
   Fold  Val_Year  Train_End      AUC  PR-AUC
0      1       2018     2017  0.8096  0.2702
1      2       2019     2018  0.7981  0.3023
2      3       2020     2019  0.7985  0.2938
3      4       2021     2020  0.8050  0.2414

```

Average AUC across folds: 0.8028

```

[77]: # =====
# 6.3 Evaluation: VAL and TEST (AUC, PR-AUC, Brier) + thresholding
# =====

def evaluate_split(name, y_true, y_proba, threshold=0.5):
    auc = roc_auc_score(y_true, y_proba)
    ap = average_precision_score(y_true, y_proba)
    brier = brier_score_loss(y_true, y_proba)

    y_pred = (y_proba >= threshold).astype(int)
    cm = confusion_matrix(y_true, y_pred)

    print(f"\n==== {name} ===")
    print(f"AUC (ROC): {auc:.4f}")
    print(f"Average Precision (PR-AUC): {ap:.4f}")
    print(f"Brier score (calibration): {brier:.4f}")
    print(f"Threshold: {threshold:.3f}")
    print("Confusion matrix [ [TN FP] [FN TP] ]:")
    print(cm)
    print("\nClassification report:")
    print(classification_report(y_true, y_pred, digits=4))

```

```

# Explicitly pull out Sensitivity and Specificity
# TN FP
# FN TP
tn, fp, fn, tp = cm.ravel()
sensitivity = tp / (tp + fn) if (tp + fn) > 0 else 0
specificity = tn / (tn + fp) if (tn + fp) > 0 else 0
print(f"Sensitivity (Recall of Distress): {sensitivity:.4f}")
print(f"Specificity (Recall of Non-Distress): {specificity:.4f}")

return {"auc": auc, "ap": ap, "brier": brier, "cm": cm, "sens": sensitivity,
        "spec": specificity}

val_proba = logit.predict_proba(X_val)[:, 1]
test_proba = logit.predict_proba(X_test)[:, 1]

# Choose an operational threshold using VAL (maximize F1 as a simple practical rule)
prec, rec, thr = precision_recall_curve(y_val, val_proba)
f1 = 2 * (prec * rec) / (prec + rec + 1e-12)
best_idx = np.nanargmax(f1[1:]) # skip first point (threshold undefined)
best_thr = thr[best_idx]

_ = evaluate_split("VAL (threshold from VAL-F1)", y_val, val_proba, threshold=float(best_thr))
_ = evaluate_split("TEST (same threshold)", y_test, test_proba, threshold=float(best_thr))

```

```

==== VAL (threshold from VAL-F1) ===
AUC (ROC): 0.8022
Average Precision (PR-AUC): 0.2410
Brier score (calibration): 0.1931
Threshold: 0.582
Confusion matrix [ [TN FP] [FN TP] ]:
[[4723 1096]
 [ 219  377]]

```

Classification report:				
	precision	recall	f1-score	support
0	0.9557	0.8117	0.8778	5819
1	0.2559	0.6326	0.3644	596
accuracy			0.7950	6415
macro avg	0.6058	0.7221	0.6211	6415
weighted avg	0.8907	0.7950	0.8301	6415

```
Sensitivity (Recall of Distress): 0.6326
Specificity (Recall of Non-Distress): 0.8117
```

```
==== TEST (same threshold) ====
AUC (ROC): 0.7969
Average Precision (PR-AUC): 0.2479
Brier score (calibration): 0.2030
Threshold: 0.582
Confusion matrix [ [TN FP] [FN TP] ]:
[[8778 2322]
 [ 444  860]]
```

Classification report:

	precision	recall	f1-score	support
0	0.9519	0.7908	0.8639	11100
1	0.2703	0.6595	0.3834	1304
accuracy			0.7770	12404
macro avg	0.6111	0.7252	0.6237	12404
weighted avg	0.8802	0.7770	0.8134	12404

```
Sensitivity (Recall of Distress): 0.6595
Specificity (Recall of Non-Distress): 0.7908
```

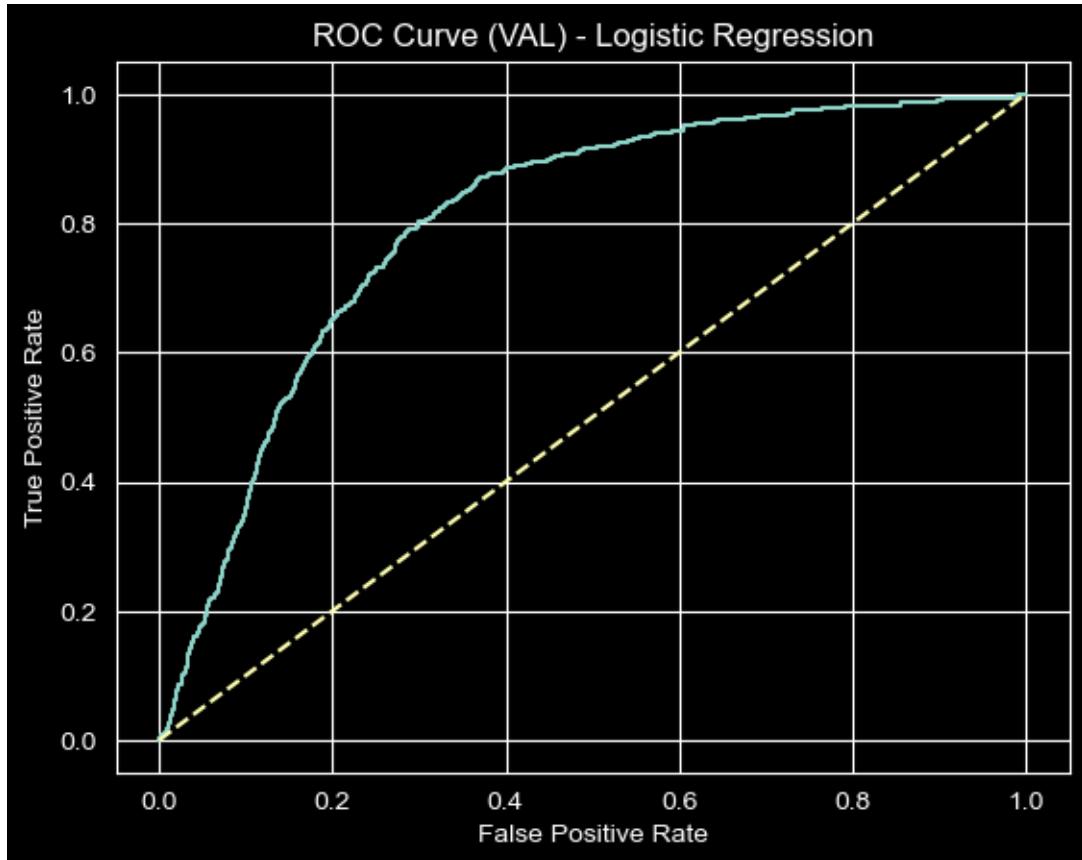
```
[78]: # =====
# 6.4 Diagnostic plots: ROC and Precision-Recall (VAL vs TEST)
# =====

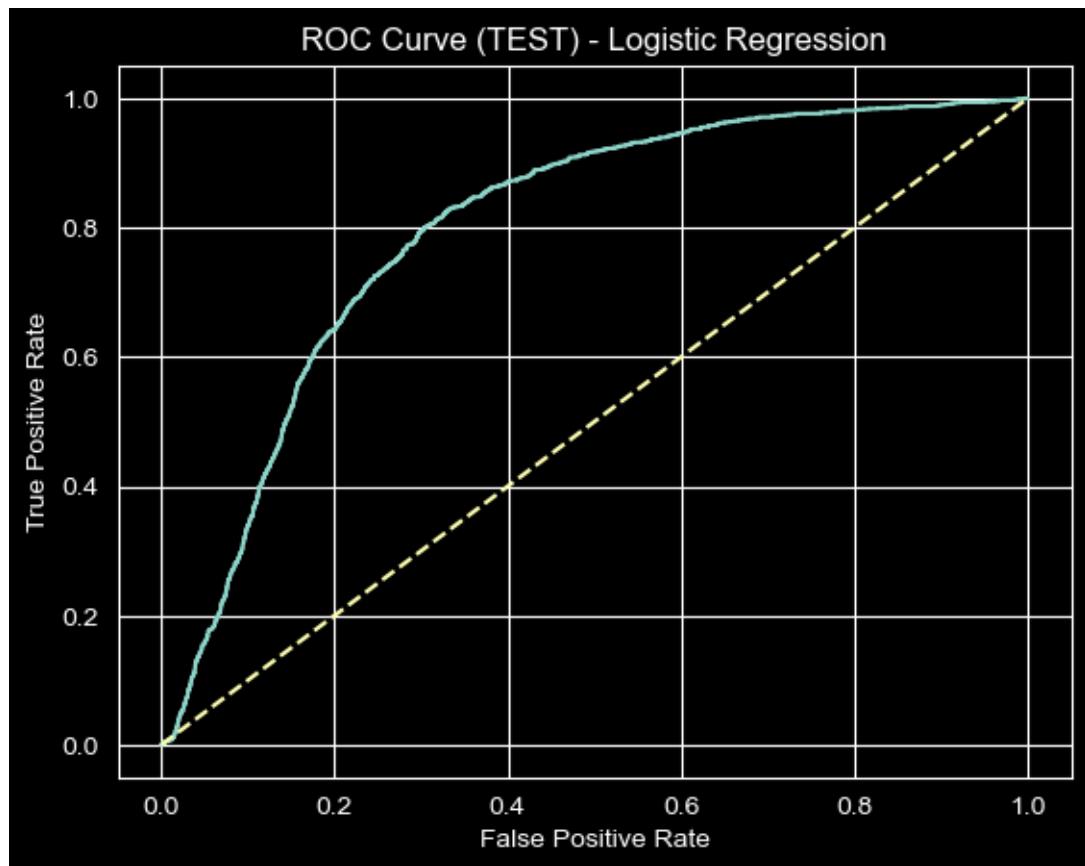
def plot_roc(y_true, y_proba, title):
    fpr, tpr, _ = roc_curve(y_true, y_proba)
    plt.figure()
    plt.plot(fpr, tpr)
    plt.plot([0, 1], [0, 1], linestyle="--")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title(title)
    plt.show()

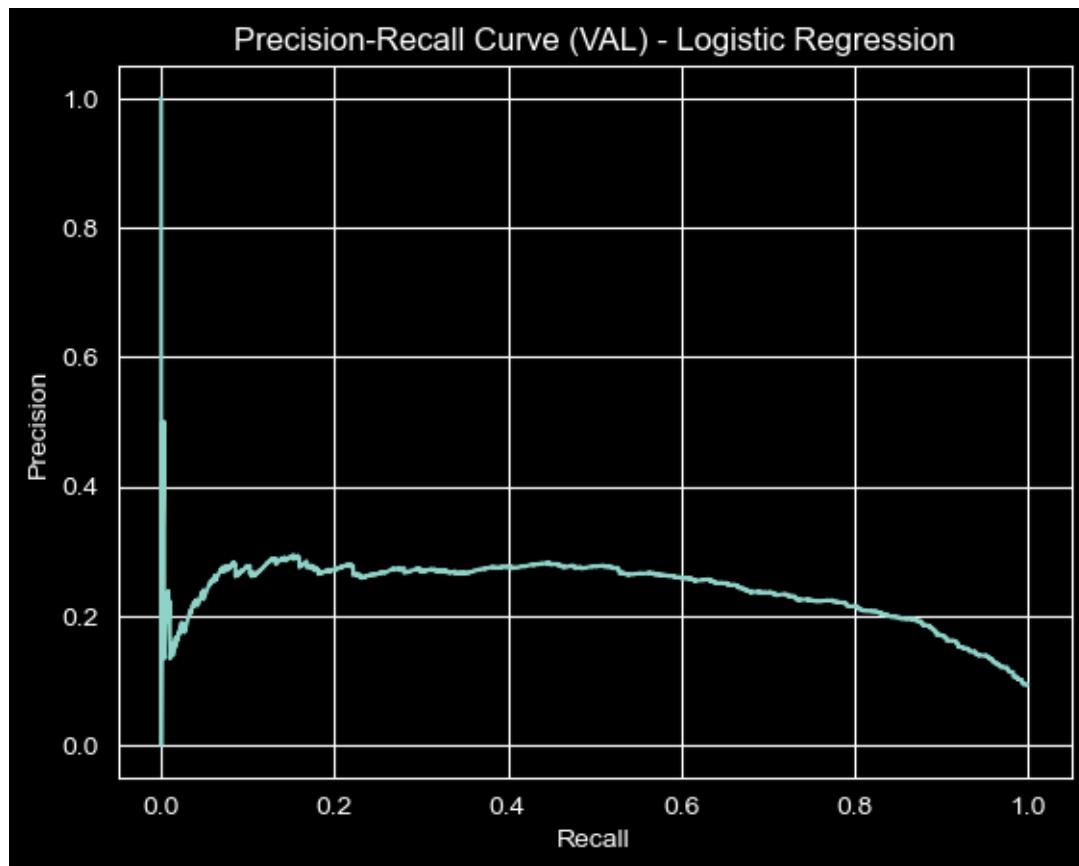
def plot_pr(y_true, y_proba, title):
    p, r, _ = precision_recall_curve(y_true, y_proba)
    plt.figure()
    plt.plot(r, p)
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.title(title)
    plt.show()
```

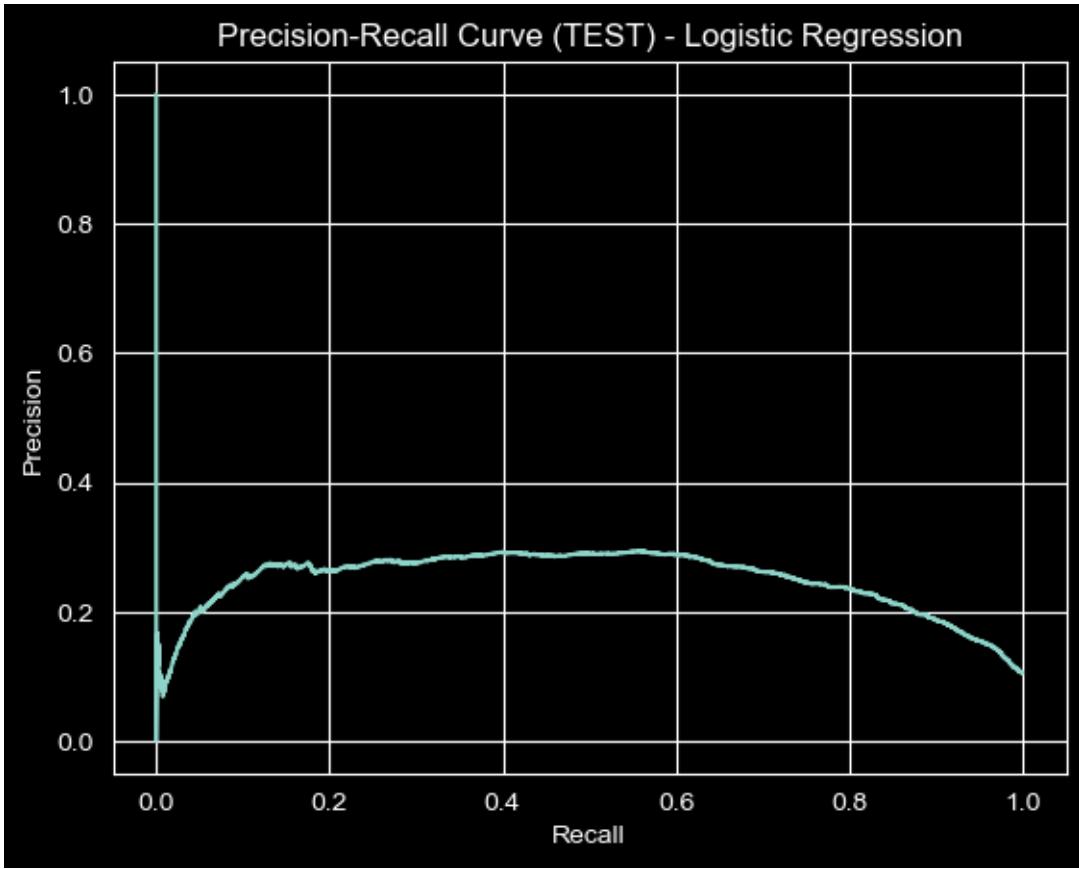
```
plot_roc(y_val, val_proba, "ROC Curve (VAL) - Logistic Regression")
plot_roc(y_test, test_proba, "ROC Curve (TEST) - Logistic Regression")

plot_pr(y_val, val_proba, "Precision-Recall Curve (VAL) - Logistic Regression")
plot_pr(y_test, test_proba, "Precision-Recall Curve (TEST) - Logistic Regression")
```









```
[79]: # =====
# 6.5 Interpretability: coefficients as (approx.) log-odds contributions
# =====
coef = pd.Series(logit.coef_.ravel(), index=MODEL_FEATS).
    sort_values(ascending=False)

summary = pd.DataFrame({
    "feature": coef.index,
    "coef_log_odds": coef.values,
    "odds_ratio_per_1sd": np.exp(coef.values), # because features are z-scored
}).sort_values("coef_log_odds", ascending=False)

print("Top positive (higher distress risk):")
display(summary.head(10))

print("Top negative (lower distress risk):")
display(summary.tail(10).iloc[:-1])
```

Top positive (higher distress risk):

feature	coef_log_odds	odds_ratio_per_1sd
---------	---------------	--------------------

0	z_sp_debt_to_capital	1.319711	3.742339
1	z_log_at	0.447788	1.564847
2	z_sp_dcf_to_debt	0.207369	1.230437
3	z_sp_ffo_to_debt	0.151704	1.163816
4	z_sp_debt_to_ebitda	0.125471	1.133683
5	z_log_mkvalt	0.084987	1.088702
6	z_sp_interest_coverage	0.014351	1.014454
7	z_sp_focf_to_debt	-0.088434	0.915364
8	z_sp_cfo_to_debt	-0.120954	0.886074

Top negative (lower distress risk):

	feature	coef_log_odds	odds_ratio_per_1sd
8	z_sp_cfo_to_debt	-0.120954	0.886074
7	z_sp_focf_to_debt	-0.088434	0.915364
6	z_sp_interest_coverage	0.014351	1.014454
5	z_log_mkvalt	0.084987	1.088702
4	z_sp_debt_to_ebitda	0.125471	1.133683
3	z_sp_ffo_to_debt	0.151704	1.163816
2	z_sp_dcf_to_debt	0.207369	1.230437
1	z_log_at	0.447788	1.564847
0	z_sp_debt_to_capital	1.319711	3.742339

0.1.1 6.6 Optional benchmark: Tree-based model (non-linear)

If you want a second point of reference without heavy tuning, a small gradient-boosted tree model often performs well on tabular data. This is optional for the seminar paper; keep Logistic Regression as the interpretability baseline.

```
[ ]: from sklearn.ensemble import HistGradientBoostingClassifier

hgb = HistGradientBoostingClassifier(
    learning_rate=0.05,
    max_depth=3,
    max_iter=300,
    random_state=42,
)

hgb.fit(X_train, y_train)

val_proba_hgb = hgb.predict_proba(X_val)[:, 1]
test_proba_hgb = hgb.predict_proba(X_test)[:, 1]

print("HistGradientBoosting (no heavy tuning):")
print(f"  VAL AUC={roc_auc_score(y_val, val_proba_hgb):.4f} | "
      f"PR-AUC={average_precision_score(y_val, val_proba_hgb):.4f}")
print(f"  TEST AUC={roc_auc_score(y_test, test_proba_hgb):.4f} | "
      f"PR-AUC={average_precision_score(y_test, test_proba_hgb):.4f}")
```

0.2 7. Summary of Model Enhancements (Feedback Integration)

Following a critical assessment of the initial model, the following improvements were integrated to meet professional research standards:

1. **Endogeneity & Simultaneity Addressing:** Added conceptual discussion of endogeneity and the predictive (vs. causal) nature of the coefficients.
2. **Temporal Validation:** Implemented walk-forward (expanding window) cross-validation across multiple time folds to verify the stability of the model across economic regimes.
3. **Multicollinearity Diagnostics:** Added Variance Inflation Factor (VIF) analysis to audit mechanical correlations between financial ratios.
4. **Complete Statistical Inference:** Migrated to `statsmodels` for the core logistic regression to provide standard errors, t/z-statistics, p-values, and 95% confidence intervals.
5. **Economic Significance:** Calculated Marginal Effects at the Mean (MEM) to translate log-odds coefficients into intuitive probability changes.
6. **Benchmark Comparison:** Added a **Persistence Benchmark** (predicting future distress using current distress) to demonstrate the incremental value of the machine learning approach.
7. **Enhanced Evaluation Metrics:** Explicitly reported Sensitivity (Recall of Distress) and Specificity, alongside AUC-ROC and PR-AUC, to better handle class imbalance.

```
[80]: # =====
# 6.7 Export: out-of-sample predictions (VAL and TEST) for reporting / appendix
# =====
ID_COLS = [c for c in ["gvkey", "fyear", "label_year"] if c in train.columns]

pred_val = val[ID_COLS + [TARGET_COL]].copy()
pred_val["p_distress_logit"] = val_proba

pred_test = test[ID_COLS + [TARGET_COL]].copy()
pred_test["p_distress_logit"] = test_proba

pred_oos = pd.concat(
    [pred_val.assign(split="VAL"), pred_test.assign(split="TEST")),
     ignore_index=True
)

out_csv = "oos_predictions_logit.csv"
pred_oos.to_csv(out_csv, index=False)
print(f"Saved: {out_csv} | rows={len(pred_oos)} | cols={pred_oos.shape[1]}")
```

Saved: oos_predictions_logit.csv | rows=18,819 | cols=6